

CS 600.226: Data Structures

Michael Schatz

Aug 30, 2018

Lecture 1: Introduction & Motivation



Welcome!

Course Webpage:

<https://github.com/schatzlab/datastructures2018>

Course Discussions:

<https://piazza.com/jhu/fall2018/600226/home>

Office Hours:

Wednesday @ 2:45pm – 4pm, Malone 323
CA office hours throughout the week ☺

Programming Language:

Java with Checkstyle and JUnit
Virtual Machine (Lubuntu) or CS acct.

Accounts for Majors (CS/CE) & Minors:

If you do not already have a personal CS departmental unix account, please complete an account request form ASAP. Check "Linux Undergrad" for account type. (Note - must be declared to be eligible.)

Accounts for Others:

We will need to make accounts. Do people need them?

CS Lab access:

Students must see Steve DiBlasio, with your J-card, in Malone G61A to get CS Lab access. The CS Lab is Malone 122 and that's where course TA/CAs will be available for help.

References and Resources

Primary Texts (Recommended, not required):

- (on-line interactive) OpenDSA, JHU version
- (print) Clifford A. Shaffer, Data Structures and Algorithm Analysis (Java Version) (Edition 3.2), available on-line and through Dover Publications.
- Peter Froehlich's Lecture notes posted to Piazza

Alternate Texts:

- Sedgewick & Wayne, Algorithms: JHU Library online edition
- Weiss, Data Structures and Algorithm Analysis in Java

Other Resources:

- Google ☺
- Code examples from Intro Programming in Java (600.107) - look in the sub-directories for examples of each topic.
- algoviz.org collection of visualizations for various data structures and algorithms
- Java API -- description of classes and methods

Grading and Help

Assessments:

- Weekly Assignments: 50% Due at 11:59pm ~one week later
- Midterm: 20% In class (~Friday Oct 12)
- Final Exam: 30% During exam week (Date TBD)
- In-class: Not graded, but there to help you!

Policies:

- Percentile scores assigned relative to the highest points awarded
- Fixed cutoffs for A+(>97); A(>93); A- (>90); B+ (>87); B (>83); B- (>80); etc
- Automatic testing and grading of coding assignments using gradescope
- **Grace period:** 10% penalty for up to 1 hour late
- **Late Days:** Five (5) chances to extend the deadline by 24 hours without any penalty

WARNING: **If you submit >1 hour late and you don't have a late day left, then you will receive 0 points**

Details:

<https://github.com/schatzlab/datastructures2018/tree/master/policies>

Course Webpage

<https://github.com/schatzlab/datastructures2018>

The screenshot shows a GitHub repository page for 'schatzlab / datastructures2018'. The repository has 47 commits, 1 branch, 0 releases, and 2 contributors. The latest commit was made 8 minutes ago by 'mischatz'. The repository contains files like assignments, lectures, policies, resources, .gitignore, and README.md. Below the repository details, there is a section titled 'JHU EN.600.226: Data Structures' with contact information for Prof. Michael Schatz and Head TA Tim Kutch. The primary goal of the course is to help students understand and implement data structures.

Materials for JHU Data Structures 2018

47 commits · 1 branch · 0 releases · 2 contributors

Branch: master · New pull request · Create new file · Upload files · Find file · Clone or download

mischatz · Update README.md · Latest commit: 8 minutes ago

assignments · Update README.md · 8 days ago

lectures · Create README.md · 8 days ago

policies · Create README.md · 8 days ago

resources · Update README.md · 7 days ago

.gitignore · .gitignore · 7 days ago

README.md · Update README.md · 8 minutes ago

README.md

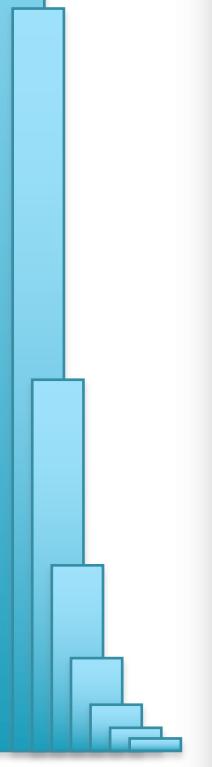
JHU EN.600.226: Data Structures

Prof: Michael Schatz (mischatz@cs.jhu.edu)
Head TA: Tim Kutch (tkutch1@jhu.edu)
Class Hours: Monday + Wednesday + Friday @ 1:30p - 2:45p in Mudd 26
Schatz Office Hours: Wednesday @ 3-4p in Malone 323 and by appointment

The primary goal of the course is for students to be comfortable in theory and leave the course empowered to understand and implement a

Course Webpage

<https://github.com/schatzlab/datastructures2018>



Schedule				
#	Date	Lecture	Readings & Resources	Assignment
1.	Th 8/30	Introduction		Sign Up for Piazza
2.	Fr 8/31	Interfaces		
	Mon 9/3	Labor Day - No class		
3.	Wed 9/5	Arrays, Generics, and Exceptions		
4.	Fri 9/7	More Arrays		HW 1 Assigned
5.	Mon 9/10	Lists		
6.	Wed 9/12	Iterators		
7.	Fri 9/14	Junit and Complexity Analysis		HW 2 Assigned
8.	Mon 9/17	Sorting		
9.	Wed 9/19	Stacks		
10.	Fri 9/21	Stacks and Queues		HW3 Assigned
11.	Mon 9/24	Stacks, Queues, and Deques		
12.	Wed 9/26	Lists		
13.	Fri 9/28	More Lists		HW4 Assigned
14.	Mon 10/1	Trees		
15.	Wed 10/3	More Trees		
16.	Fri 10/5	Graphs		
17.	Mon 10/8	Midterm Review 1		
18.	Wed 10/10	Midterm Review 2		
19.	Fri 10/12	Midterm!		
20.	Mon 10/15	Graph Searching		
21.	Wed 10/17	Graphs		HW5 Assignment

Piazza

<https://piazza.com/jhu/fall2018/600226/home>

The screenshot shows a web browser window for the Piazza platform. The URL in the address bar is <https://piazza.com/jhu/fall2018/600226/home>. The page title is "Michael Schatz - Fall 2018". The main content area displays a note titled "Welcome to Piazza!" by Michael Schatz. The note text reads:

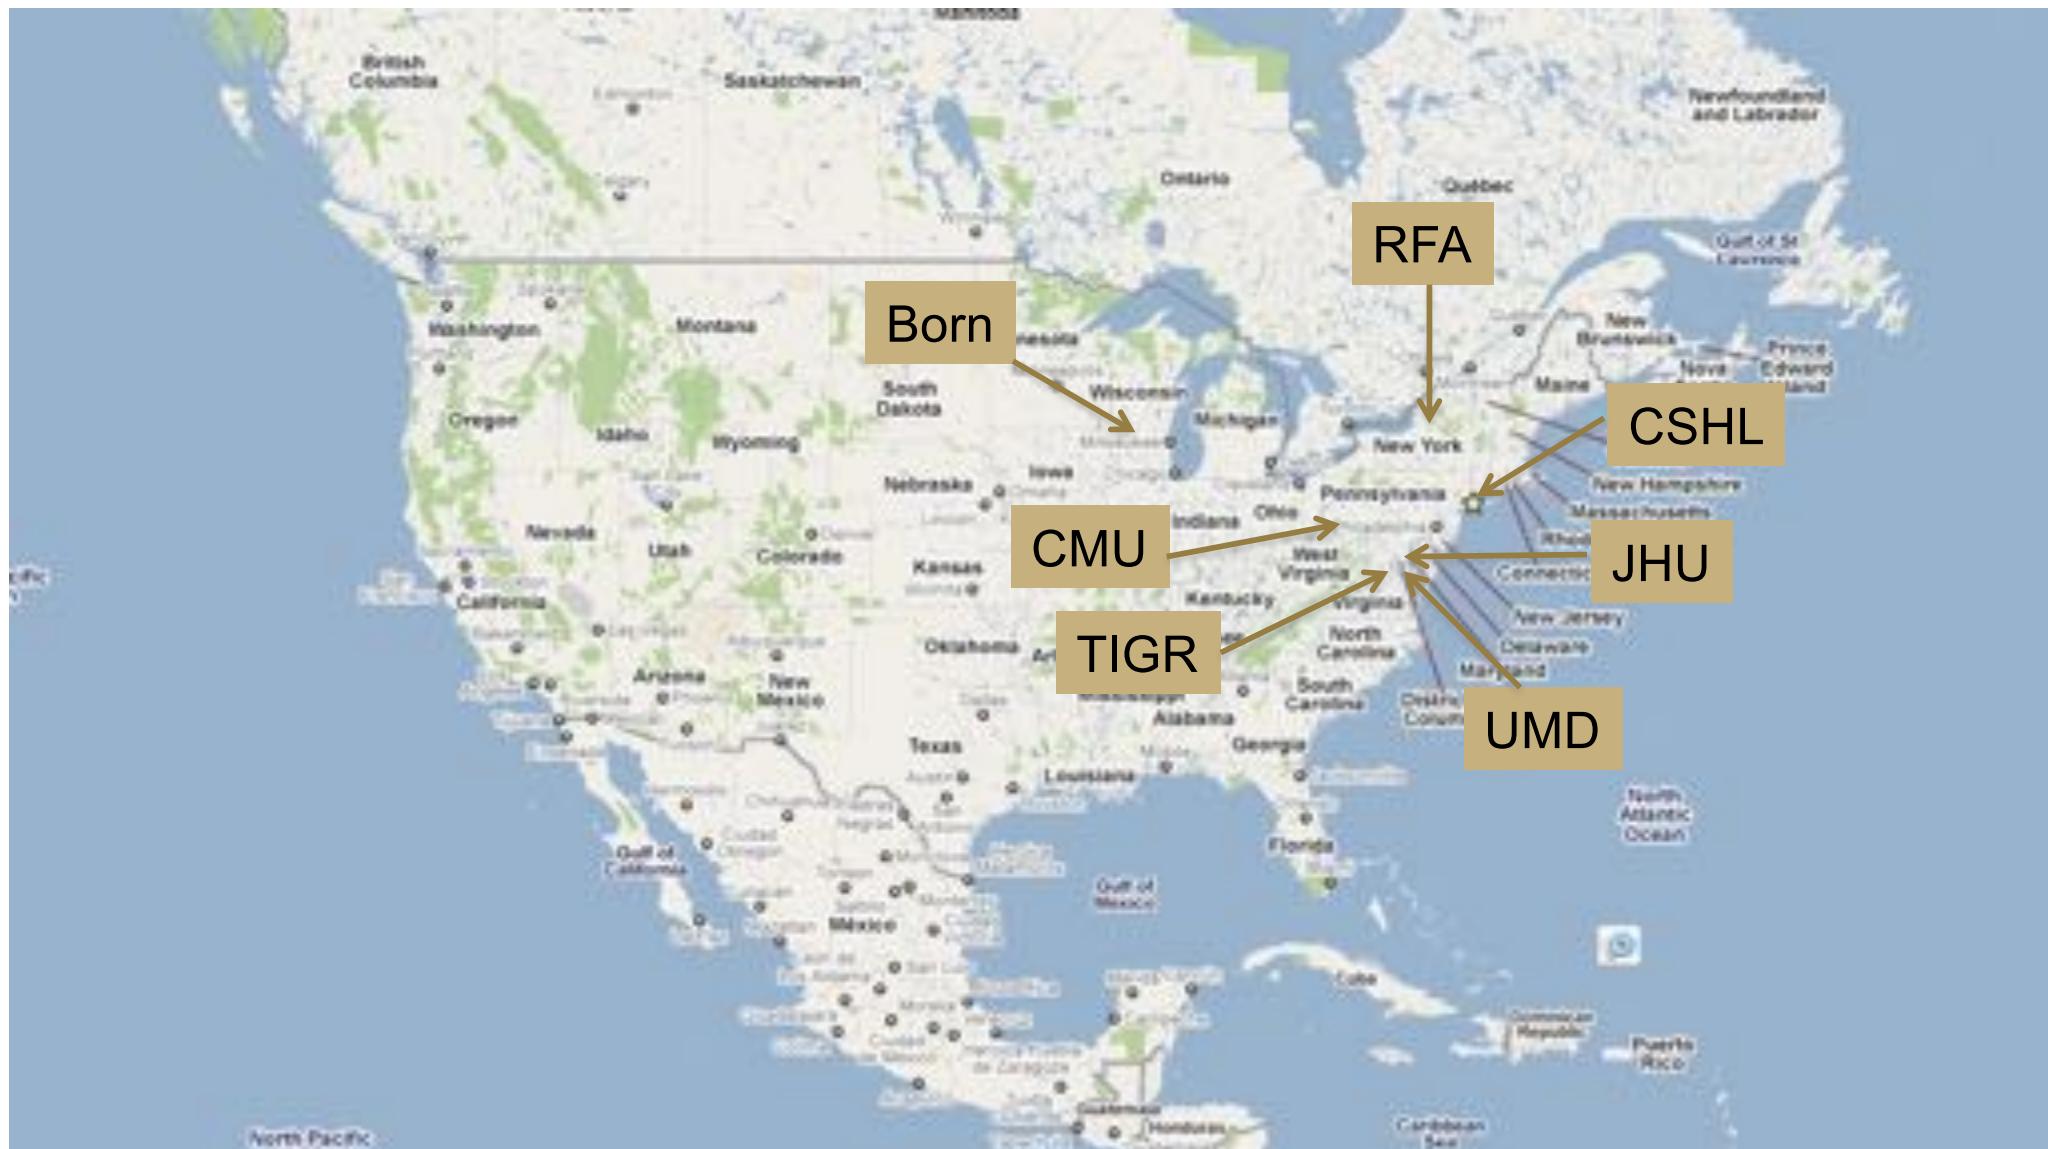
Welcome to Piazza! We'll be conducting all class-related discussion here this term. The quicker you begin asking questions on Piazza (rather than via email!), the quicker you'll benefit from the collective knowledge of your classmates and instructors. We encourage you to ask questions when you're struggling to understand a concept — you can even do so anonymously.

Below the note, there is a "good note" button and a link to "View history". On the left sidebar, under "FAVORITES", there is a post titled "Welcome to Piazza!". The post text is:

Welcome to Piazza!
Piazza is a Q&A platform designed to get you great answers from classmates and instructors fast. We've put together the

On the right side of the note, there is a "view following" button and a "2 views" counter. At the bottom of the note, it says "Updated 8 days ago by Michael Schatz".

A Little About Me



Schatzlab Overview



Human Genetics

Role of mutations in disease

Nattestad et al. (2018)
Feigin et al. (2017)



Agricultural Genomics

Genomes & Transcriptomes

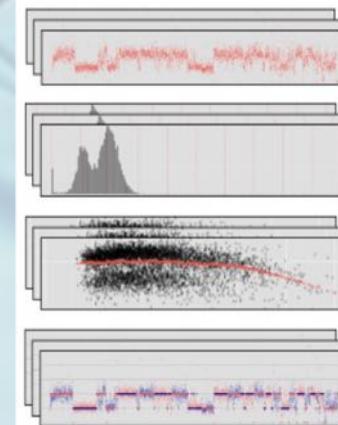
Lemmon et al. (2016)
Ming et al. (2015)



Algorithmics & Systems Research

Ultra-large scale biocomputing

Stevens et al. (2015)
Marcus et al. (2014)



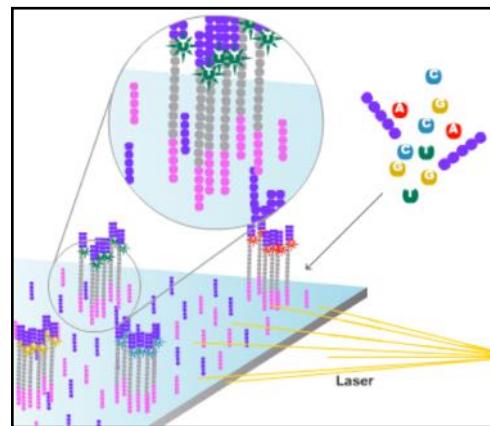
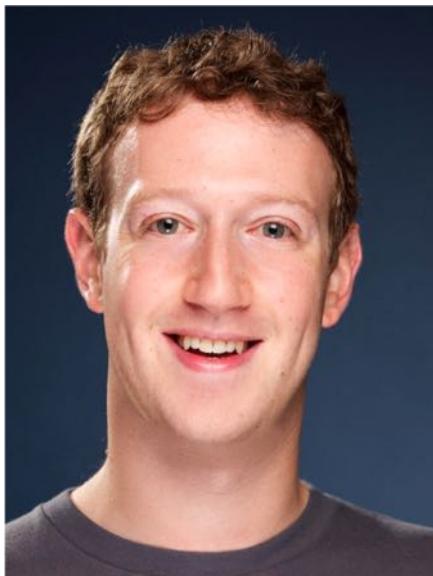
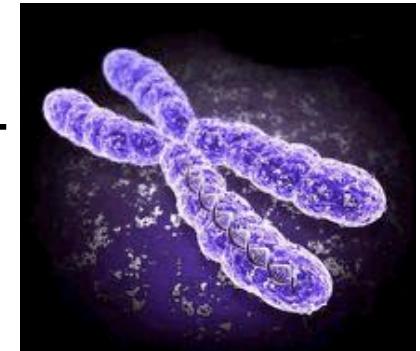
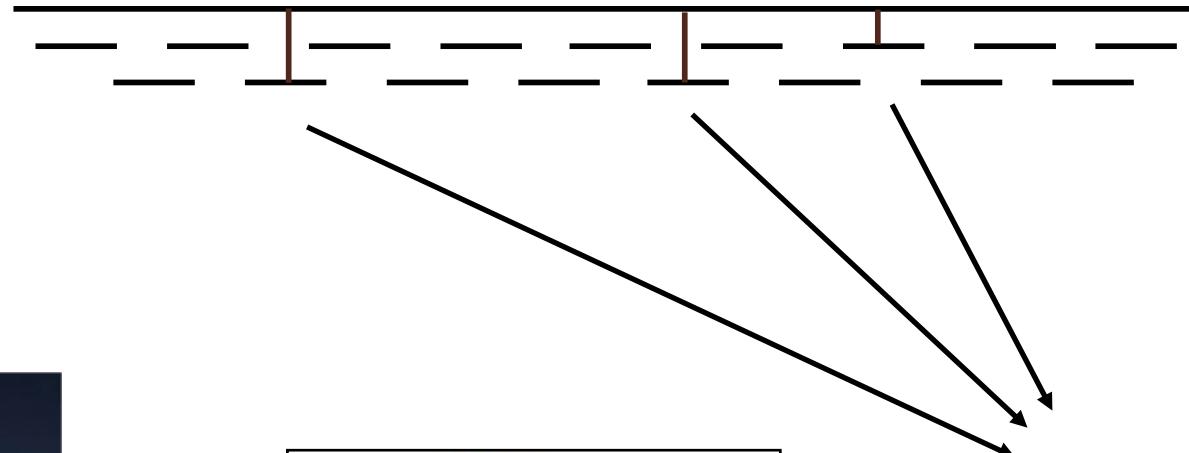
Single Cell & Single Molecule

CNVs, SVs, & Cell Phylogenetics

Sedlazeck et al. (2018)
Garvin et al. (2015)

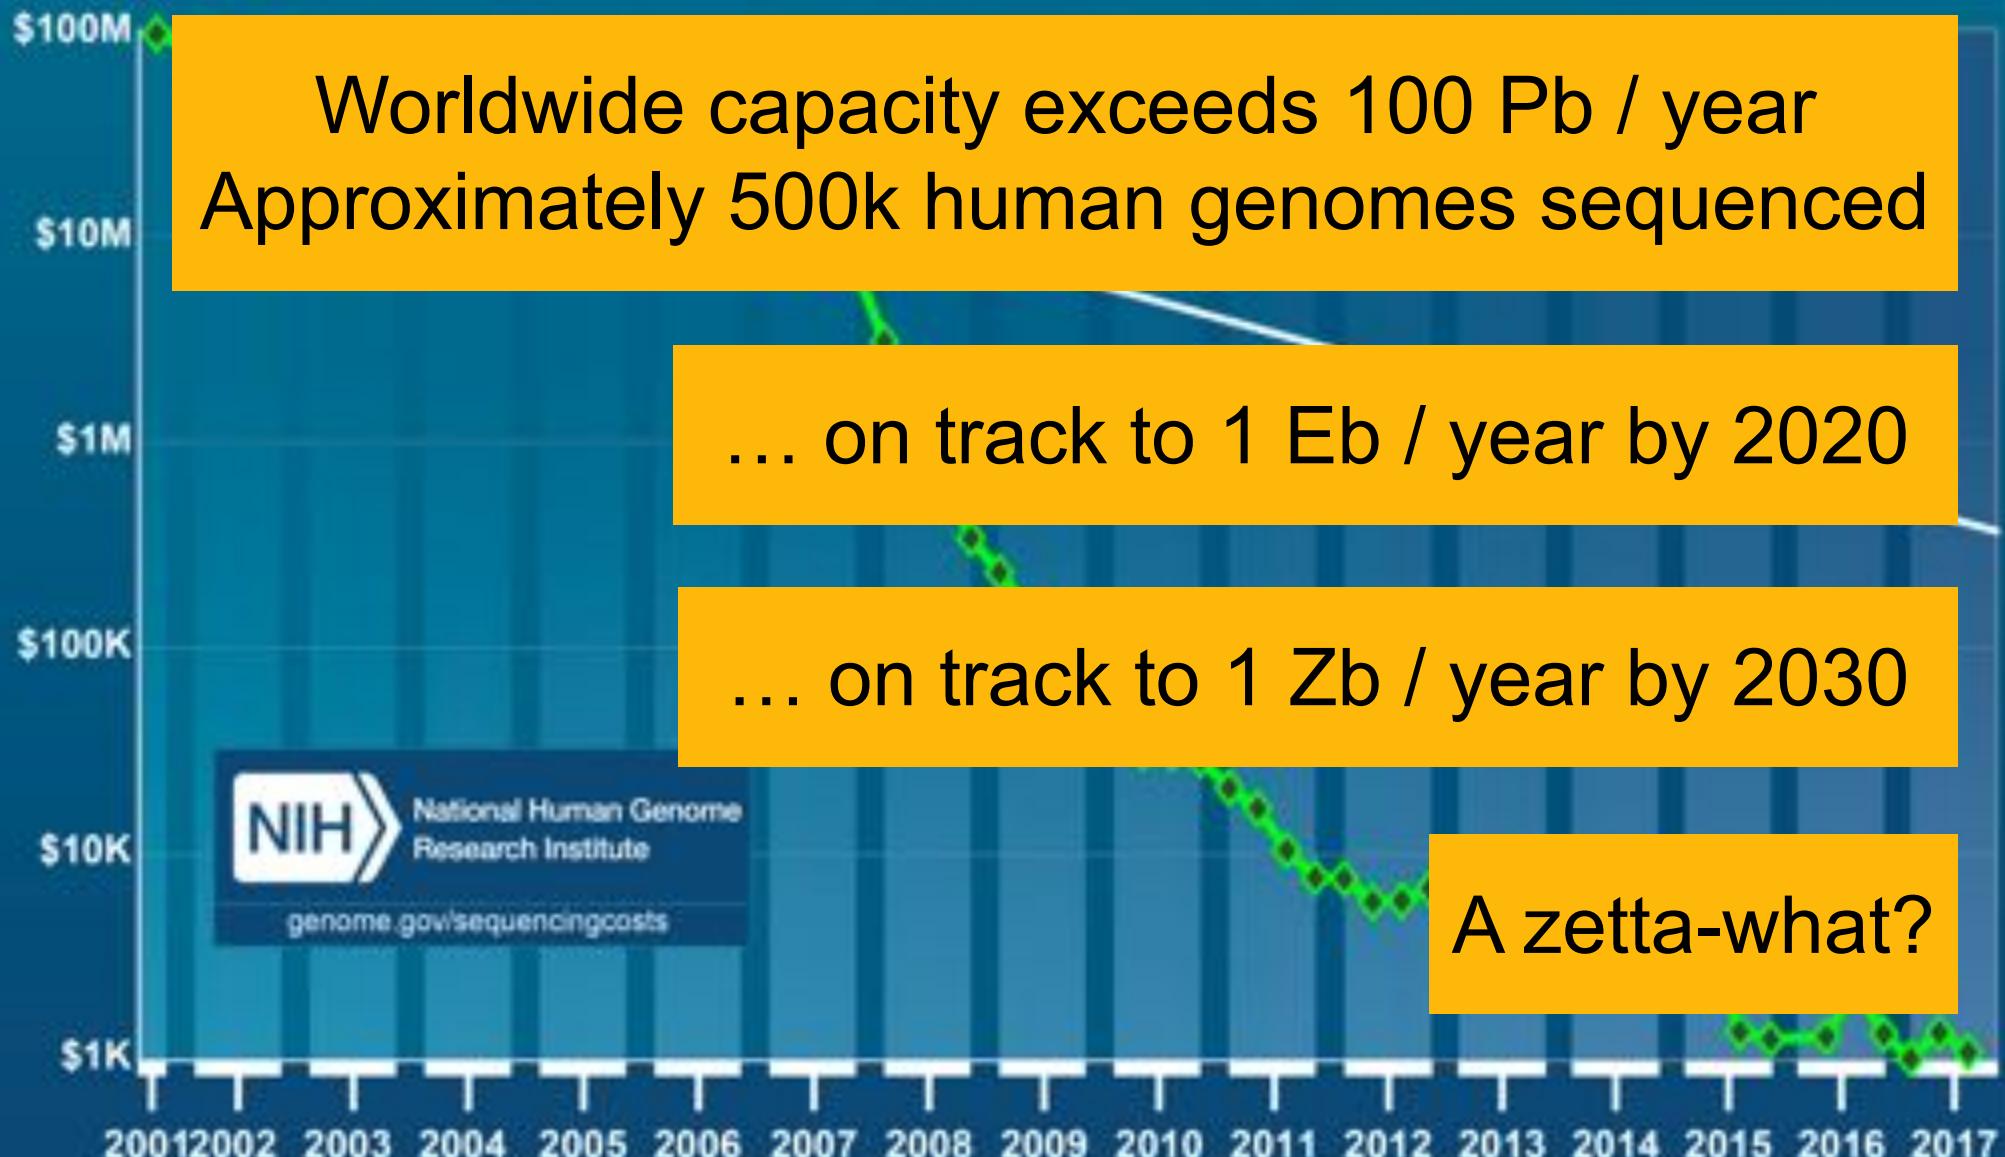
Personal Genomics

How does your genome compare to the reference?



Heart Disease _____
Cancer _____
Technology _____
Innovator _____

Cost per Genome



How much is a zettabyte?

Unit	Size	$\sim 2^x$
Byte	1	2^0
Kilobyte	1,000	2^{10}
Megabyte	1,000,000	2^{20}
Gigabyte	1,000,000,000	2^{30}
Terabyte	1,000,000,000,000	2^{40}
Petabyte	1,000,000,000,000,000	2^{50}
Exabyte	1,000,000,000,000,000,000	2^{60}
Zettabyte	1,000,000,000,000,000,000,000	2^{70}

How much is a zettabyte?



100 GB / Genome
4.7GB / DVD
~20 DVDs / Genome

X

10,000,000,000 Genomes

=

1ZB Data
200,000,000,000 DVDs



150,000 miles of DVDs
~ ½ distance to moon



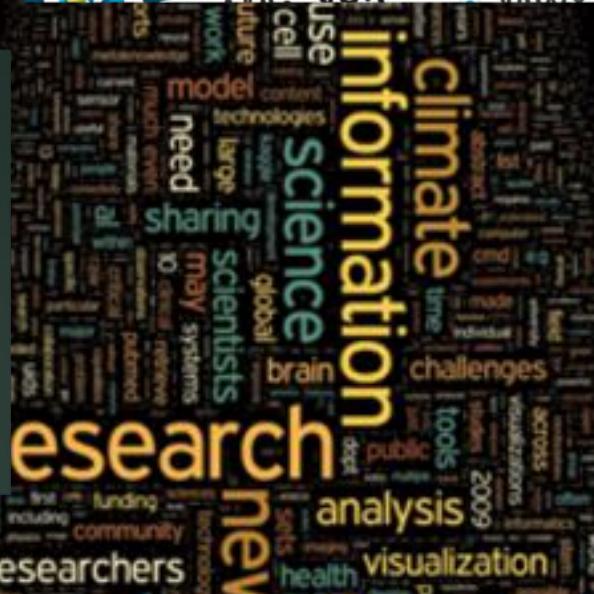
Both currently ~100PB
And growing exponentially

11 February 2013 | \$16

scientific Science



S
e

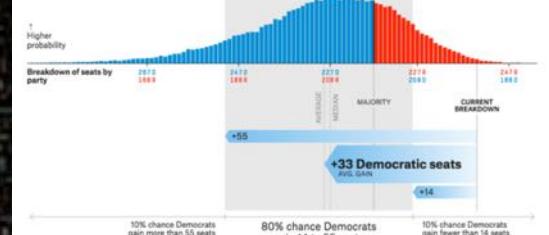


Forecasting the race for the House

Updated Aug. 20, 2018, at 9:48 PM

5 in 7

Chance Democrats win control (73.1%)

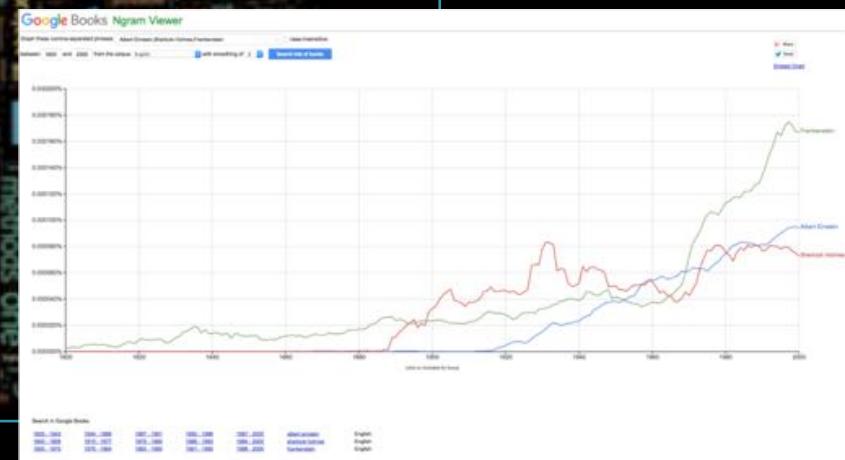


2 in 7

Chance Republicans keep control (26.9%)

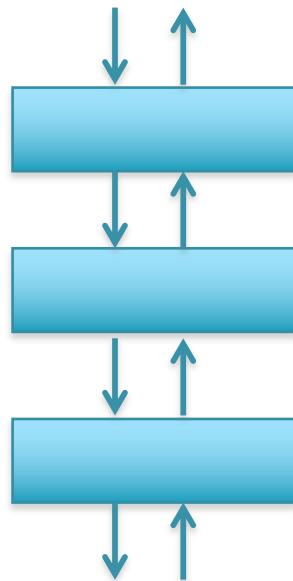


a



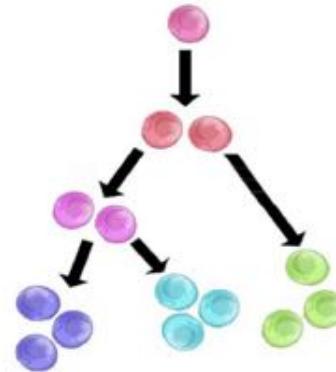
Data Structures

Lists



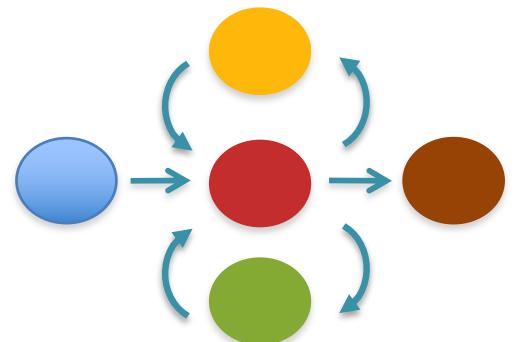
- Single/Double
- Stacks/Queues/Deques
- Skip

Trees



- Binary, AVL Trees
- Heaps
- Self Balancing

Graphs



- Graph Representations
- Traversing
- Union Find

*Building, searching, traversing, analyzing
Make you big-data superheros ☺*

Instagram

Search 5:47 PM 91%

Instagram

neil_shubin Cape Cod National Seashore



69 likes

neil_shubin A perfect combination: Glorious sunrise, beach cliffs all to myself, and coffee. Saying goodbye to... [more](#)

Add a comment...

10 HOURS AGO

Suggestions for You

See All

Home Search + Heart People

Search 5:46 PM 91%

johns hopkins

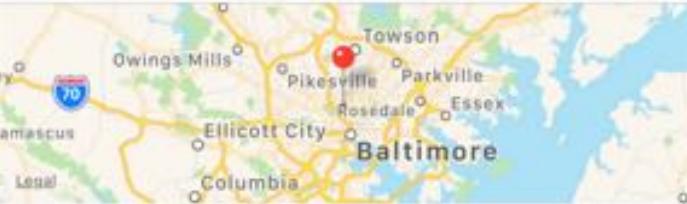
Top People Tags Places

- Johns Hopkins University
3400 N Charles St, Baltimore, MD
- Johns Hopkins Hospital
1800 Orleans St
- Johns Hopkins Children's Center
1800 Orleans St, Baltimore, MD
- Johns Hopkins Medicine
- Johns Hopkins Bloomberg School of Public Health
615 N Wolfe St, Baltimore, MD
- Johns Hopkins Bayview Medical Center
4940 Eastern Ave, Baltimore, MD
- Johns Hopkins School of Advanced International Studies
1740 Massachusetts Ave NW, Washington, DC

q w e r t y u i o p
a s d f g h j k l
z x c v b n m
123 ☺ space Search

Search 9:33 AM 100%

Johns Hopkins University



RELATED BUSINESS

johnshopkinsu Johns Hopkins University

TOP POSTS



Home Search + Heart People

Data Structures of Instagram

Incredibly popular app:

~800M active users

>20B photos, >60M per day!

<https://www.quora.com/How-many-photos-are-being-uploaded-on-Instagram-daily>

How to find all photos near a given site?

Modern clock speed: 1 instruction / nanosec

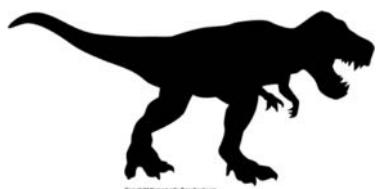
Practical processing speed: 1000 photos / sec

1M seconds = ~11.5 days

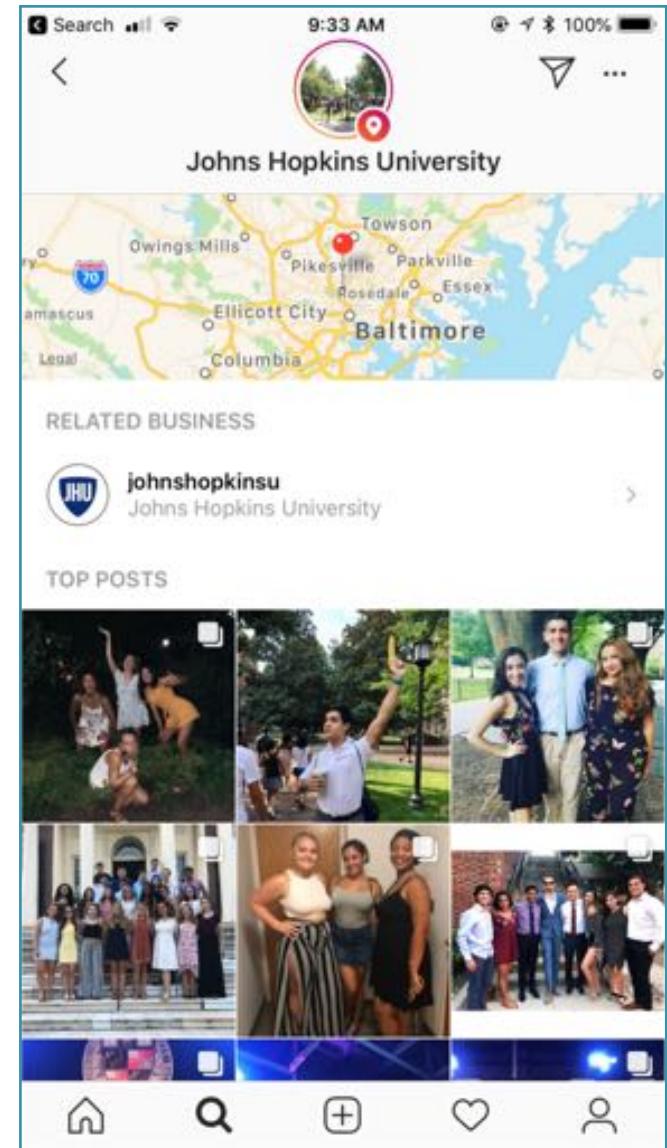
20B photos / 1000 photos / s = 20M sec
= ~230 days

What if all users search at the same time?

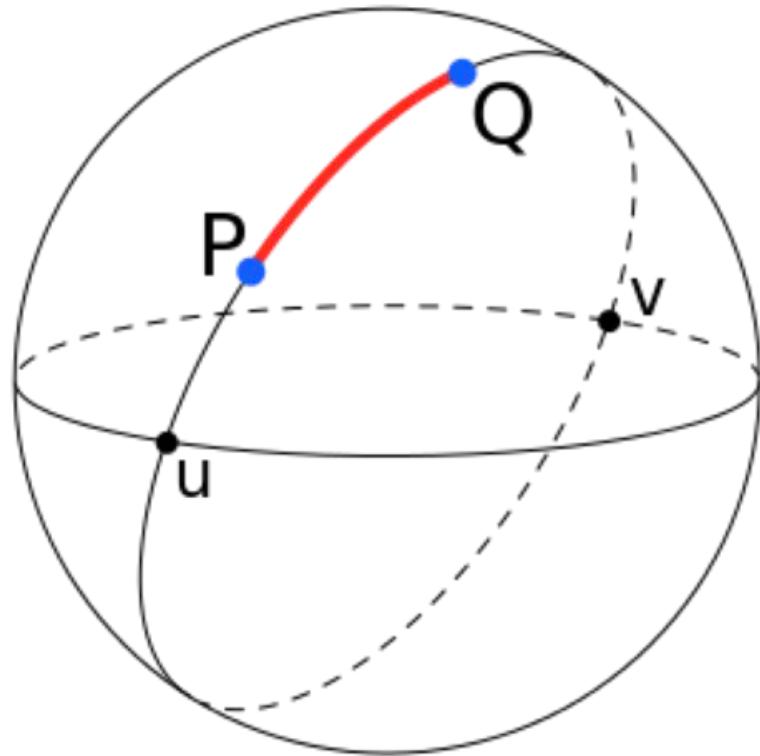
$230 \text{ days} * 800M \text{ users} = 184B \text{ days}$
 $= \sim 500M \text{ years}$



How can we make it go faster?



Data Structures of Instagram



$$\Delta\sigma = \arctan \frac{\sqrt{(\cos \phi_2 \cdot \sin(\Delta\lambda))^2 + (\cos \phi_1 \cdot \sin \phi_2 - \sin \phi_1 \cdot \cos \phi_2 \cdot \cos(\Delta\lambda))^2}}{\sin \phi_1 \cdot \sin \phi_2 + \cos \phi_1 \cdot \cos \phi_2 \cdot \cos(\Delta\lambda)}.$$

https://en.wikipedia.org/wiki/Great-circle_distance

Inside Instagram

Search: JHU

Where: 39.32N 76.62W

Photo #1

Where: 37.77N 122.41W (SFO)

URL: instagram.com/p/1

Photo #2

Where: 20.63N 76.77W (Cuba)

URL: instagram.com/p/2

...

Photo #3526224

Where: 39.32N 76.63W (JHU!)

URL: instagram.com/p/3526224

Show me the photos!

Linear Search (aka Brute force): try all 20B photos

#1: 37.77N 122.41W: No

#2: 20.63N 76.77W: No

#3: 21.30N 157.85W: No

...

#3,526,224 39.32N 76.63W: Yes!

...

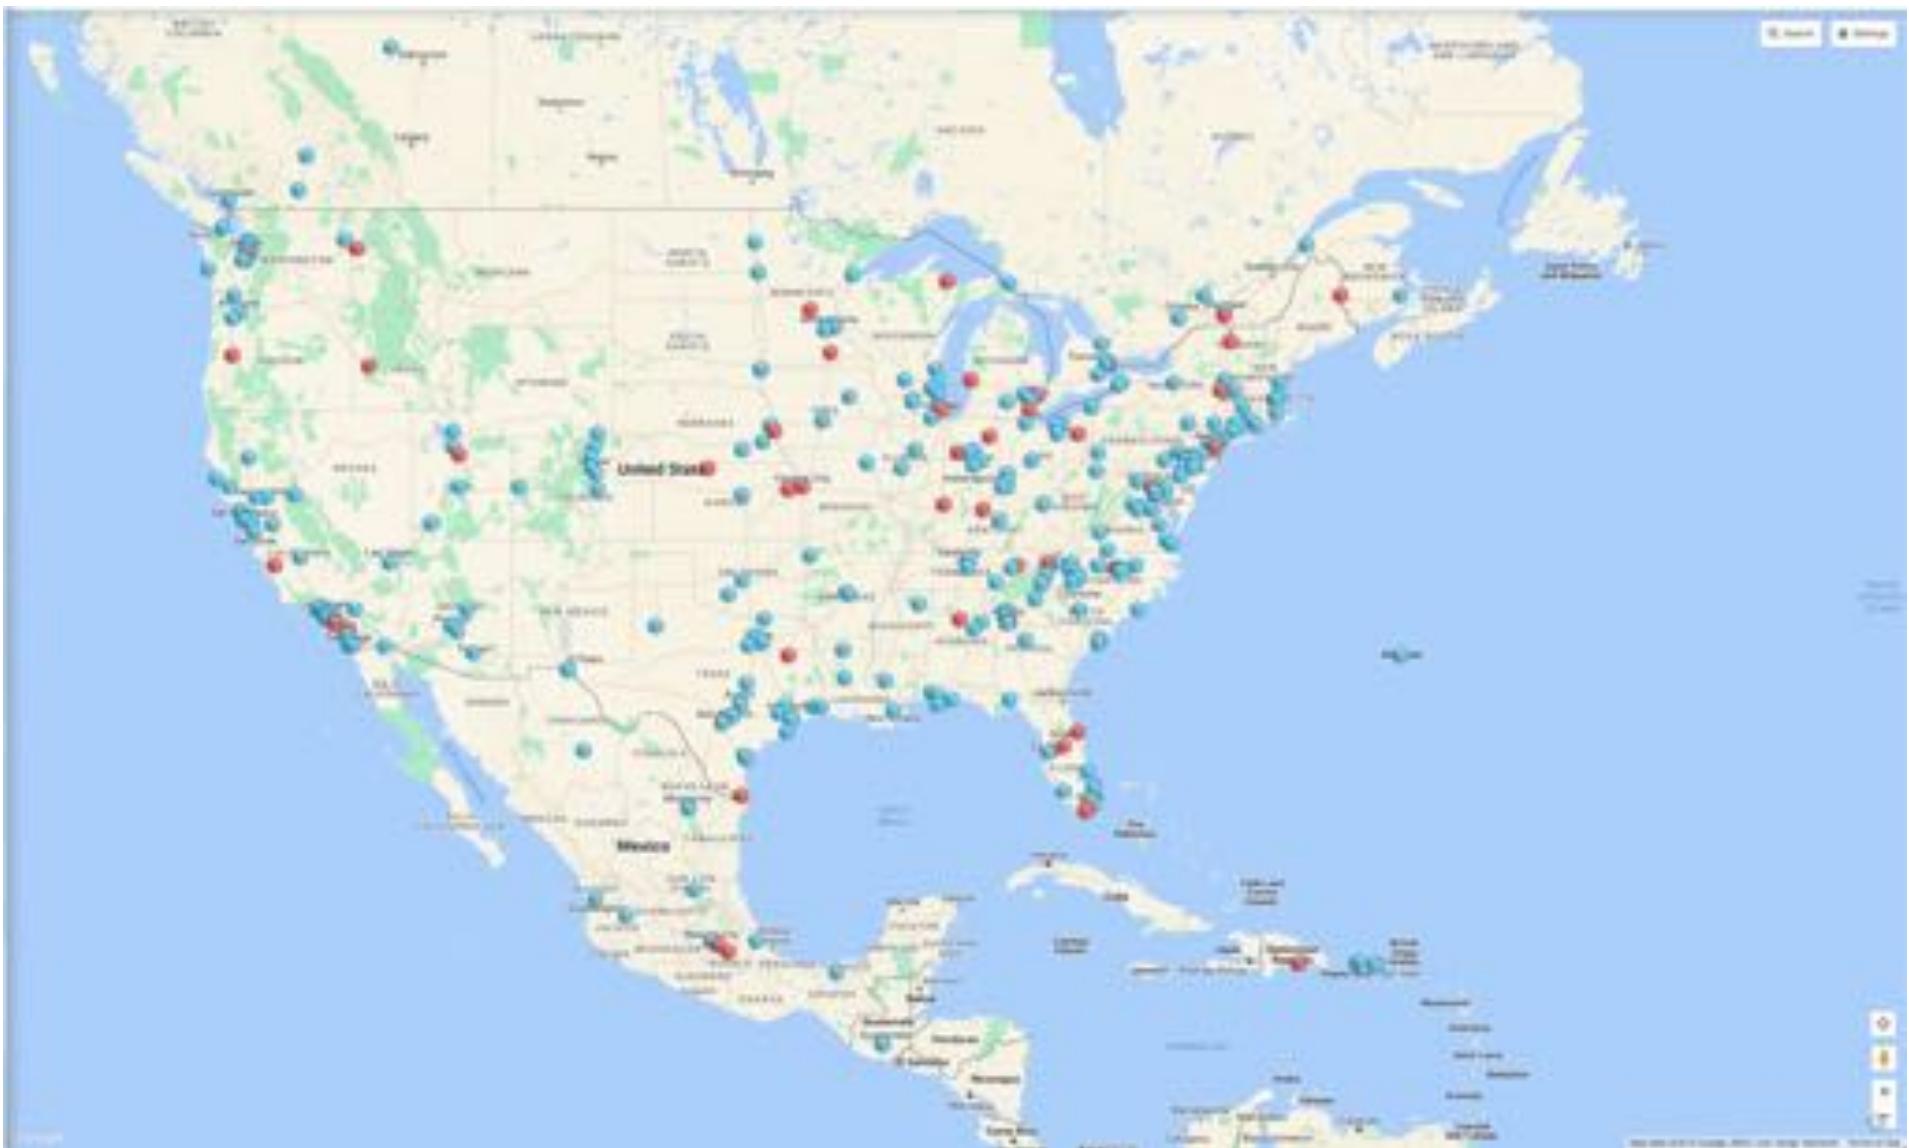
#19,999,999,999 48.85N 2.34E No

#20,000,000,000 35.65N 139.83E No

If you get really lucky you might find a few nearby photos quickly that you can return first

What happens if there are no photos at the search site?

Show me the photos!



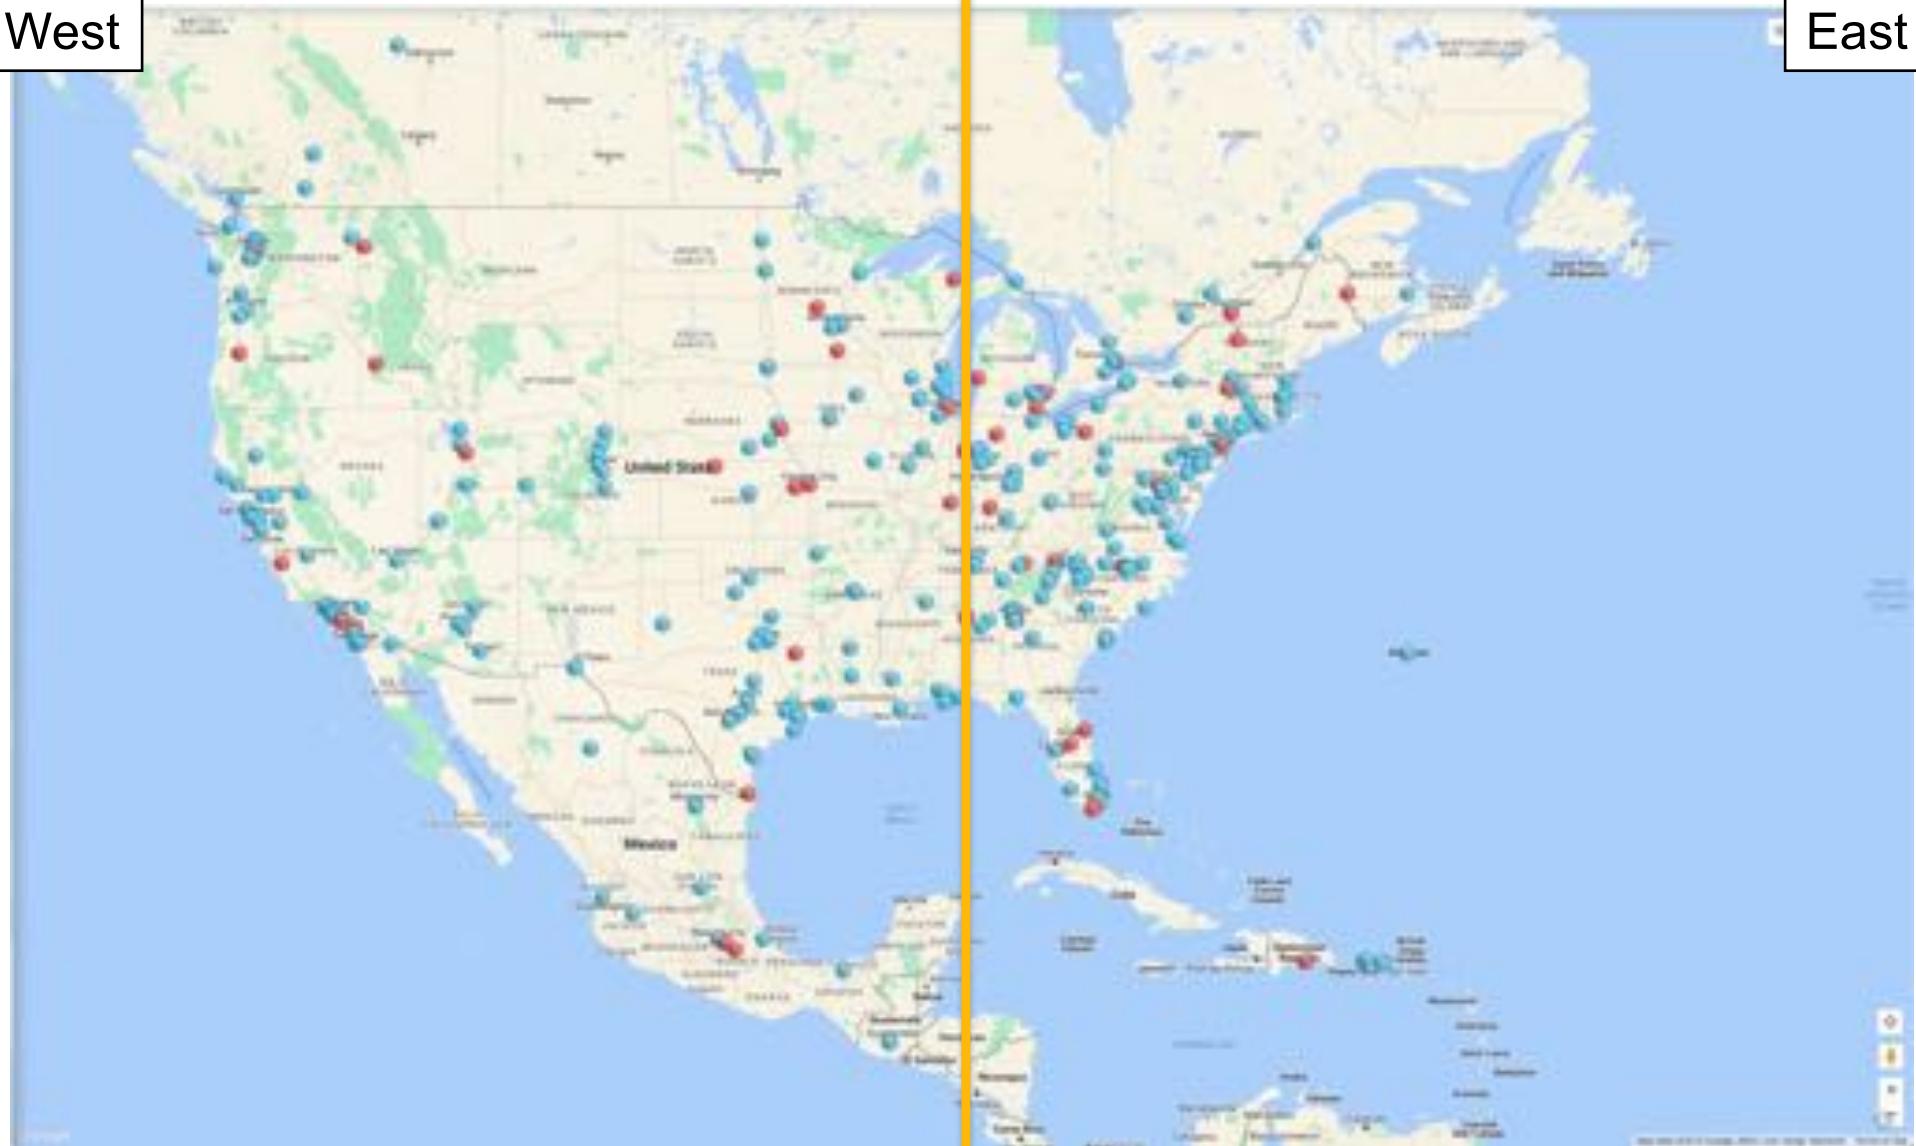
What can you do to speed up the search?

Note: The computer can only “see” one photo at a time

Show me the photos!

West

East

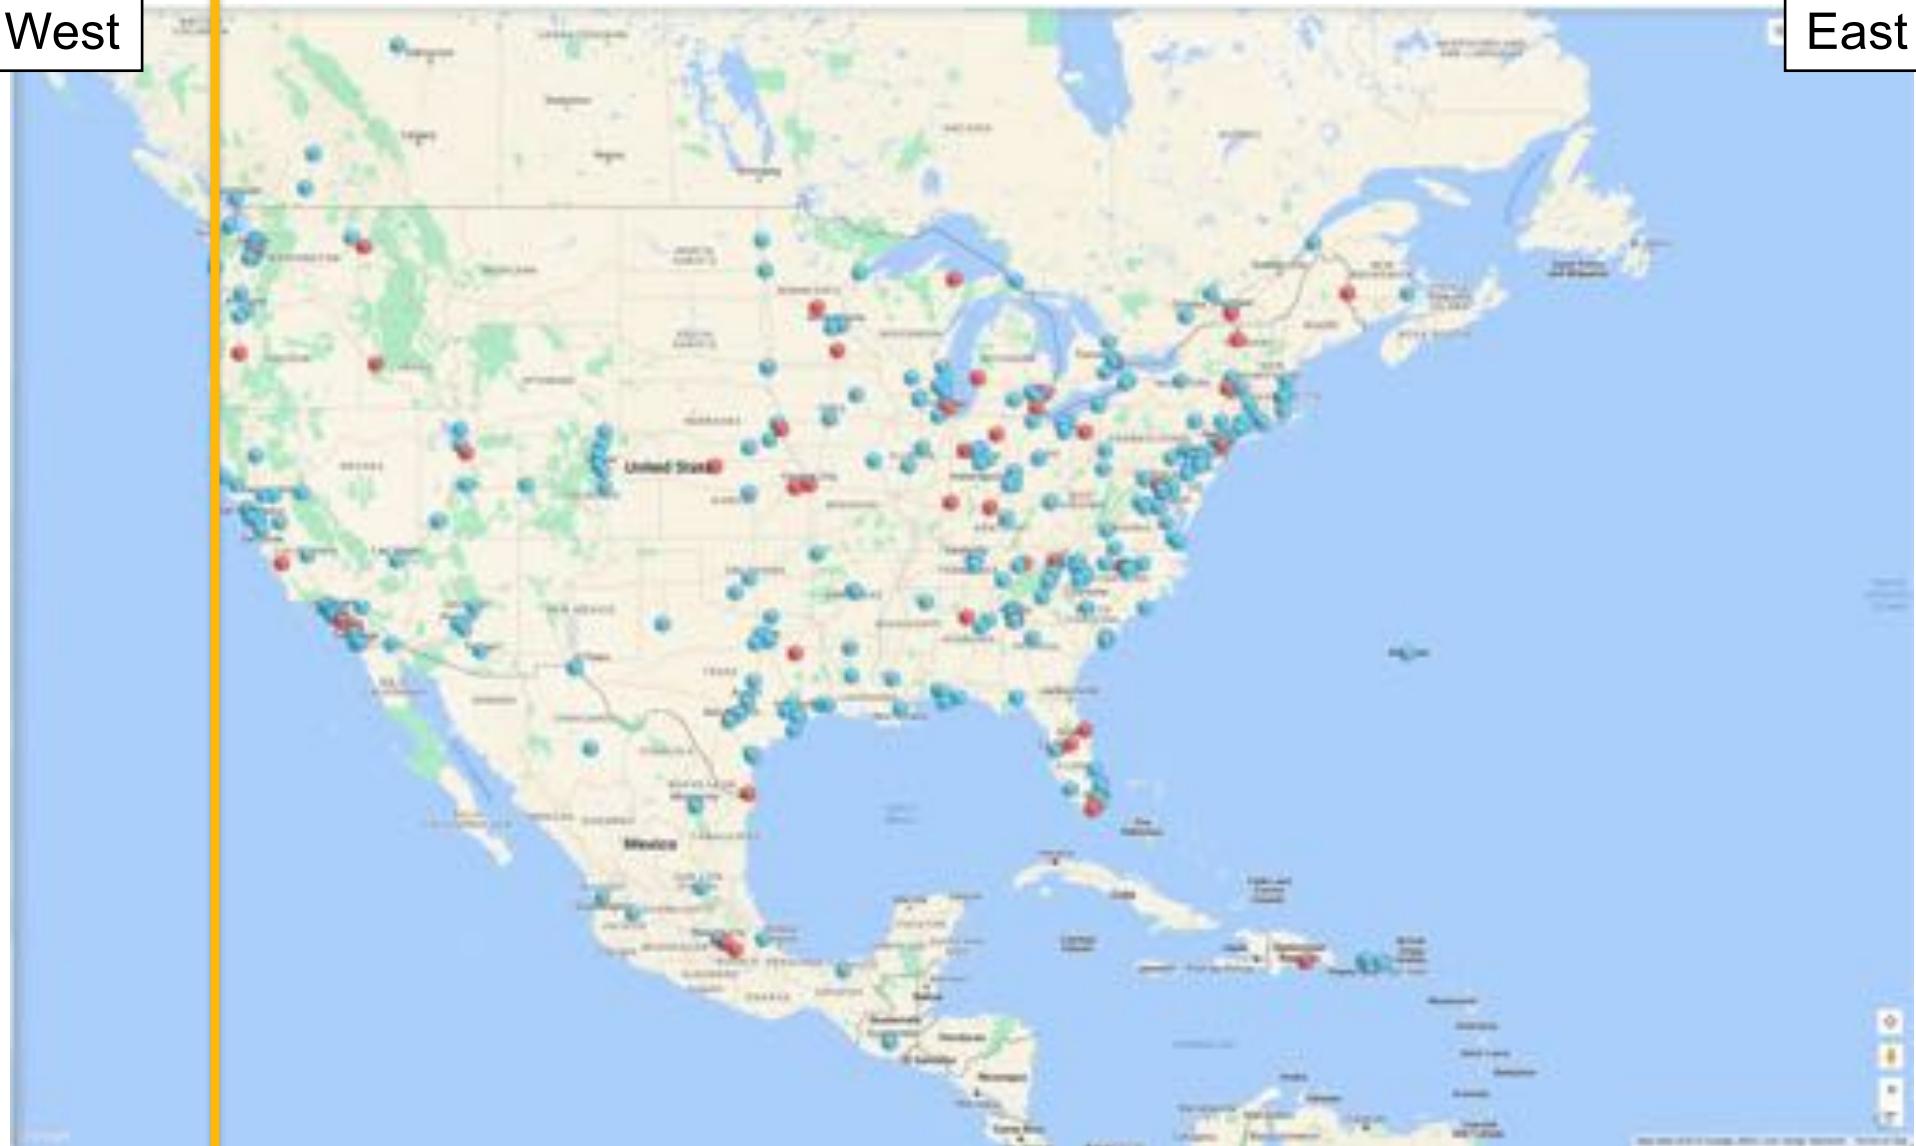


Partition the data into 2 lists, each search takes half as long!

Show me the photos!

West

East

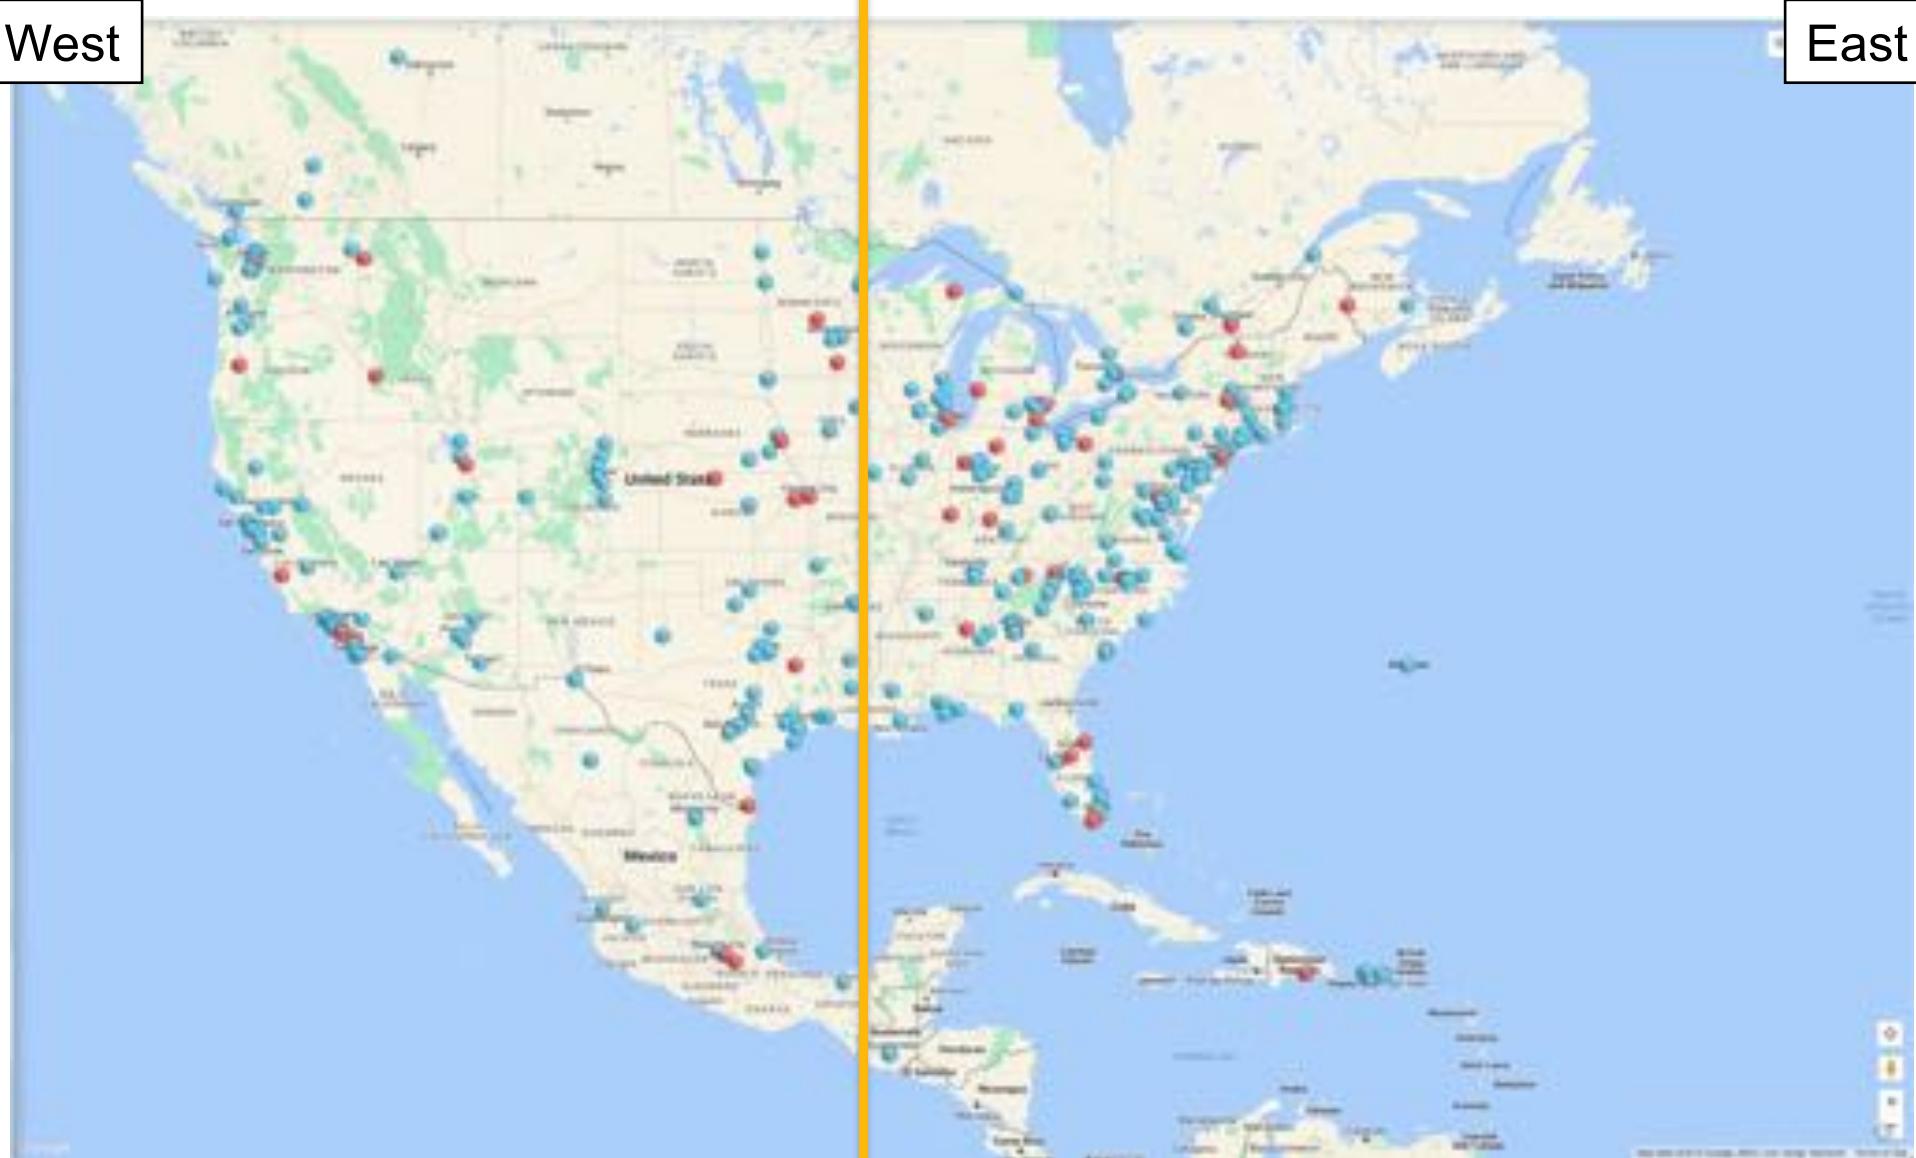


Why is this a bad split? What would be the perfect split?

Show me the photos!

West

East



Ideal split will be exactly 50/50 (median east-west coordinate of sites)

Show me the photos!

West-West



West-East



East-West



East-East



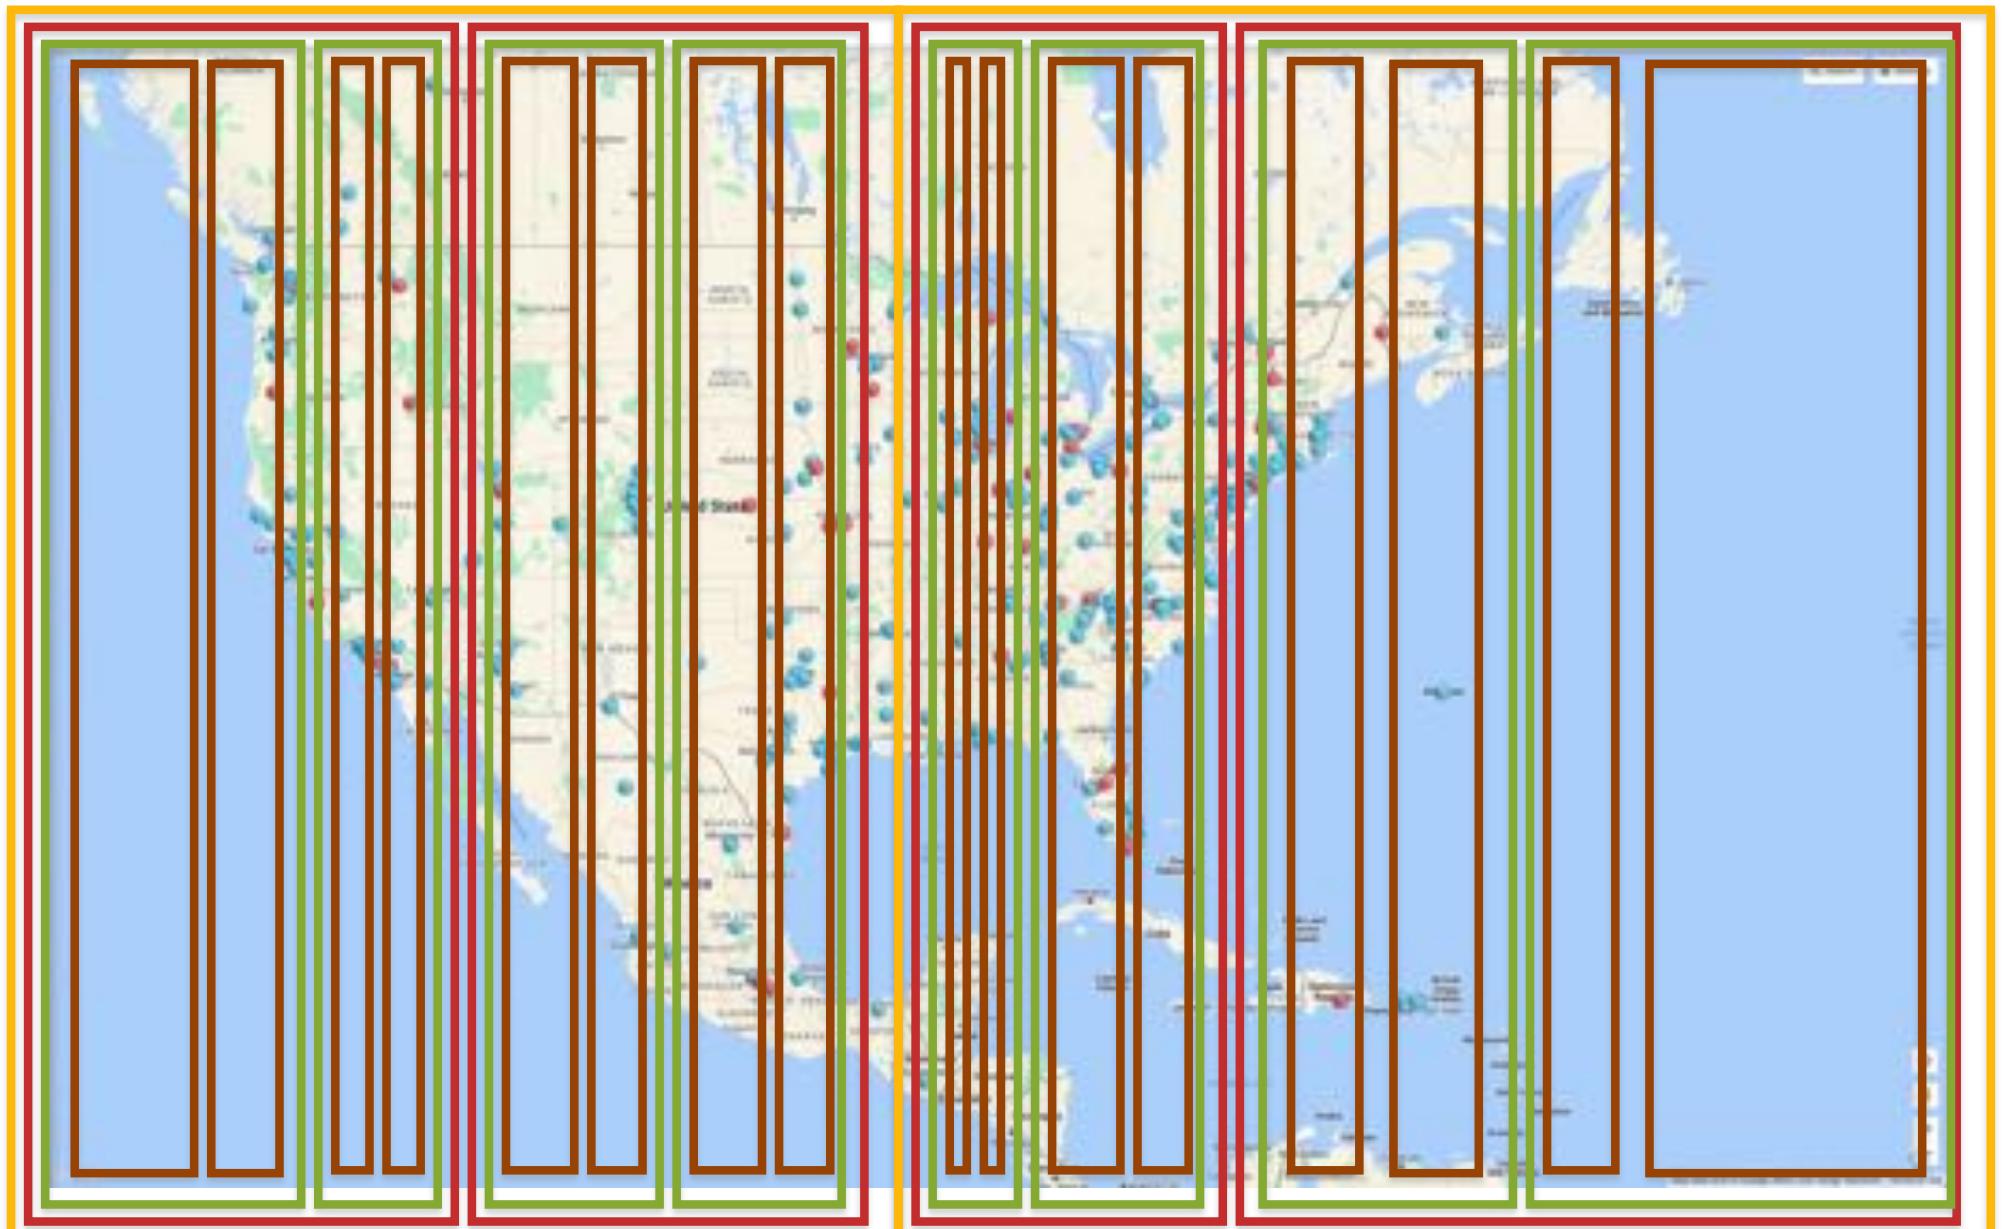
Partition again! Each sublist has $N/4$ elements!

Show me the photos!



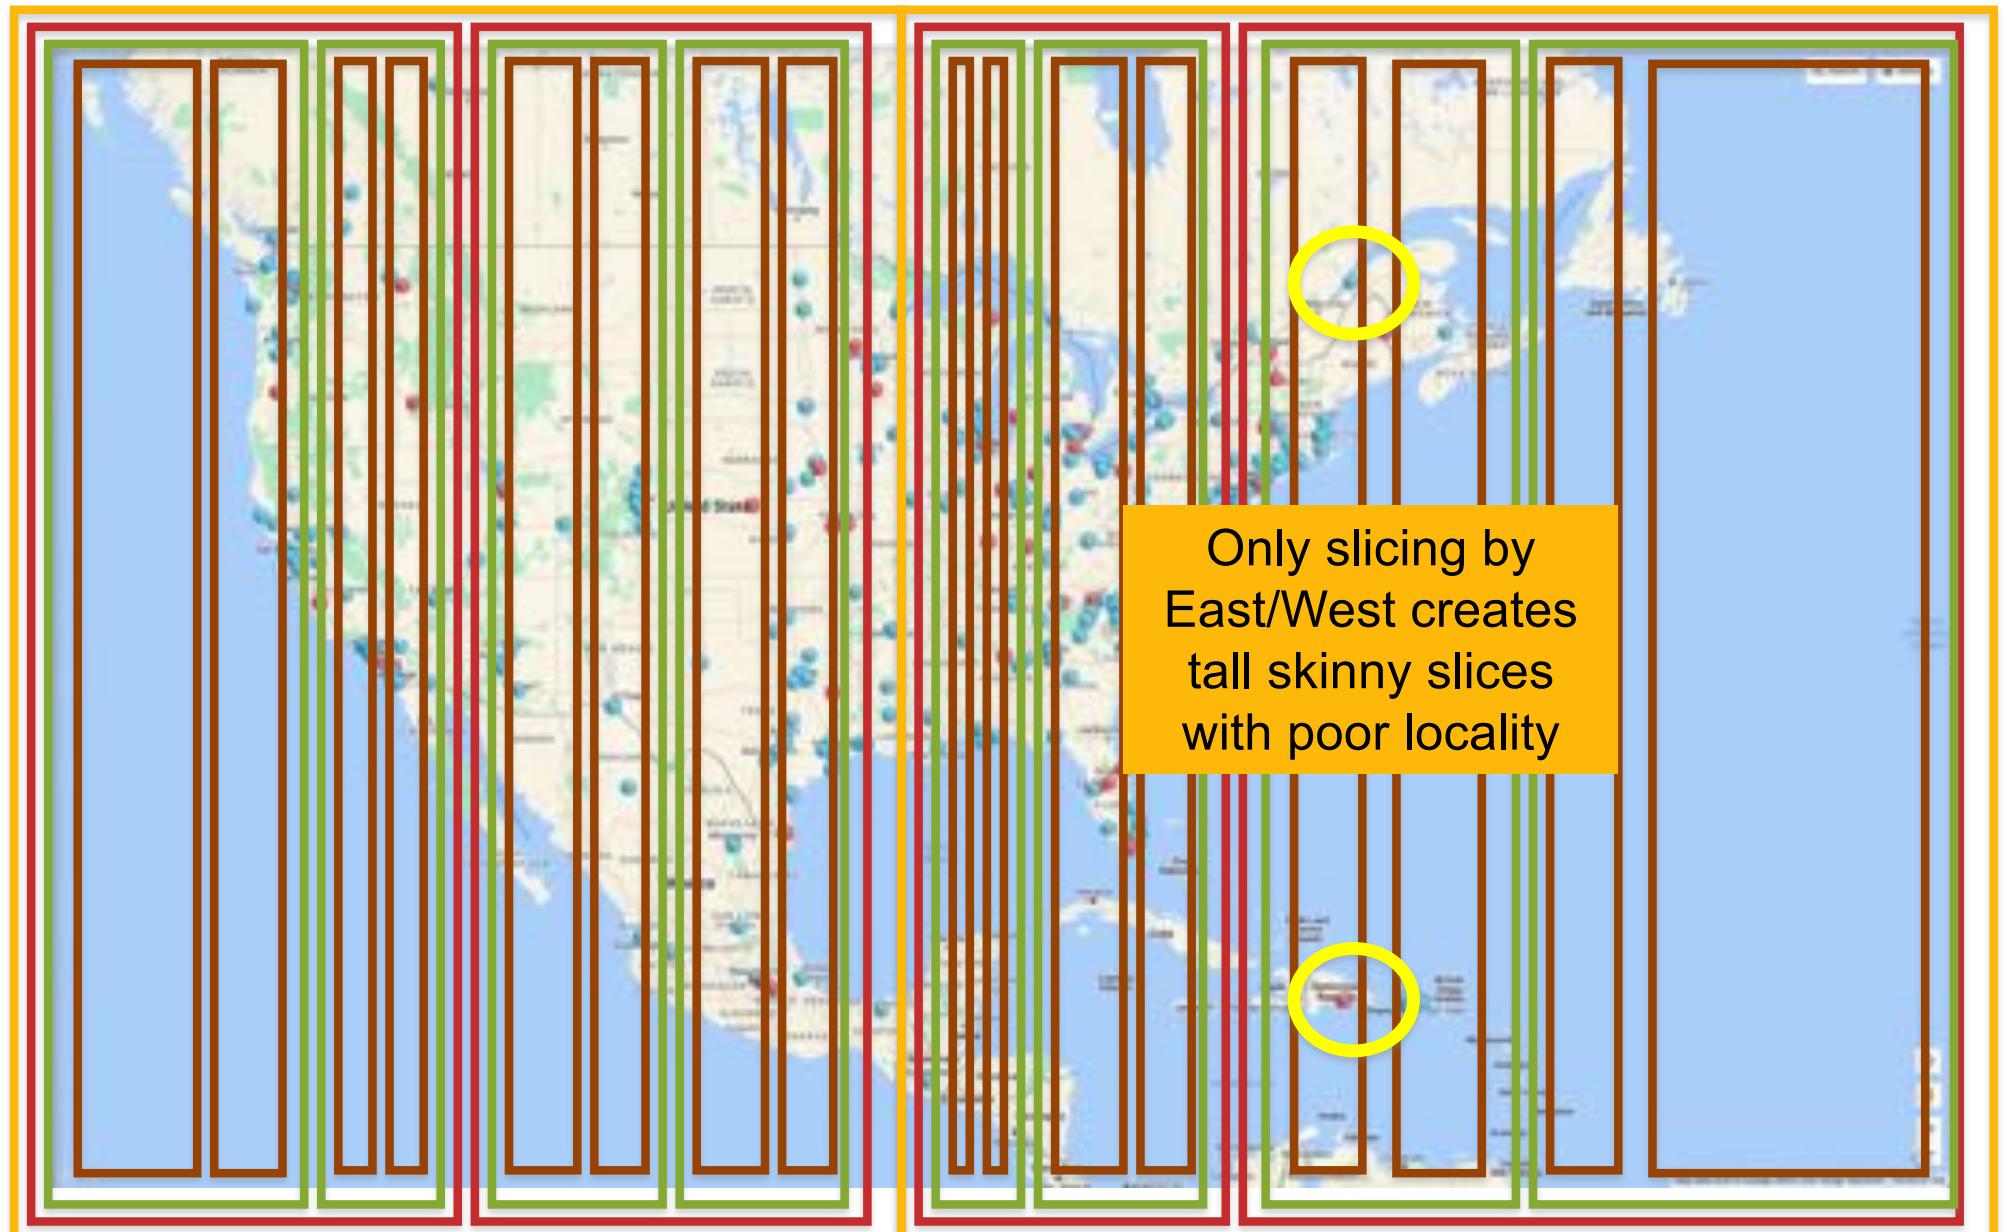
Partition again! Each sublist has $N/8$ elements!

Show me the photos!



Partition again! Each sublist has $N/16$ elements!

Show me the photos!

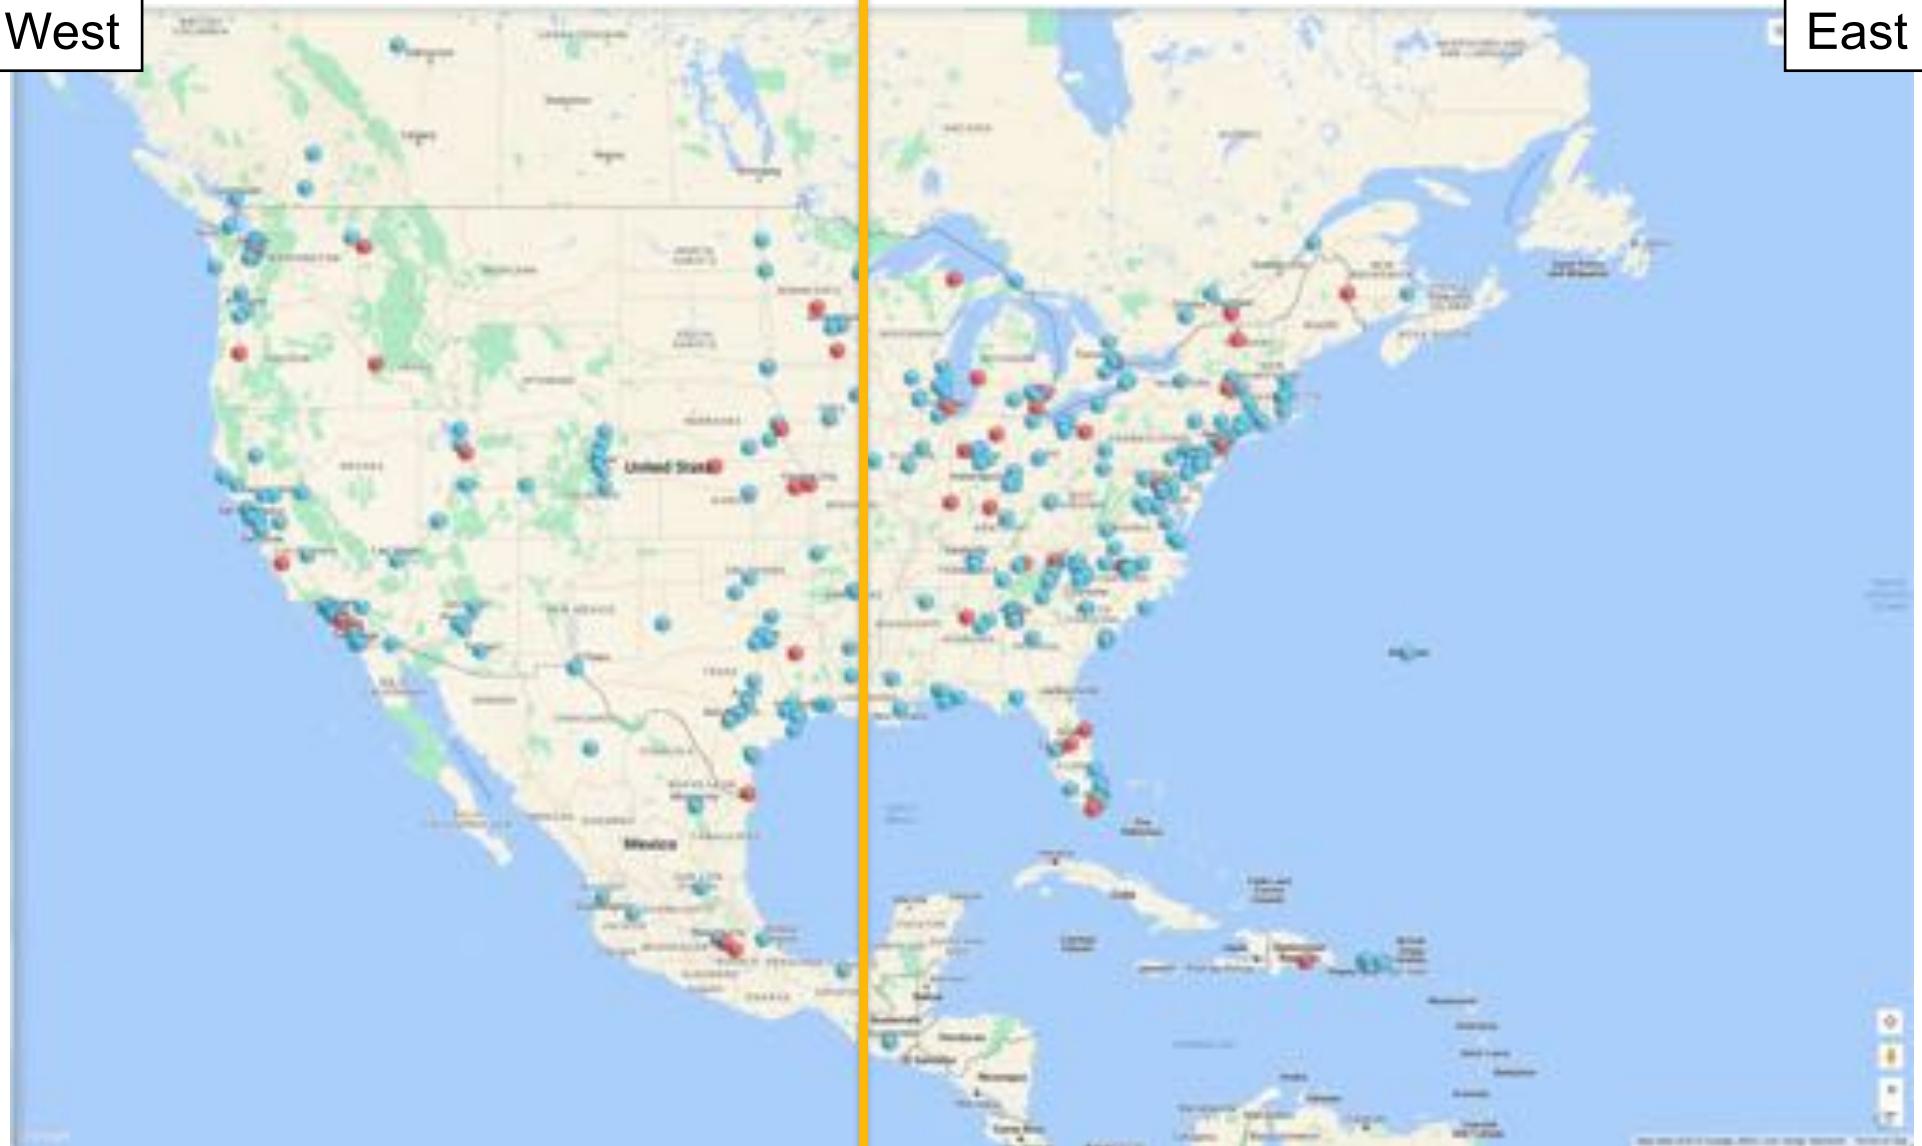


Partition again! Each sublist has $N/16$ elements!

Show me the photos!

West

East



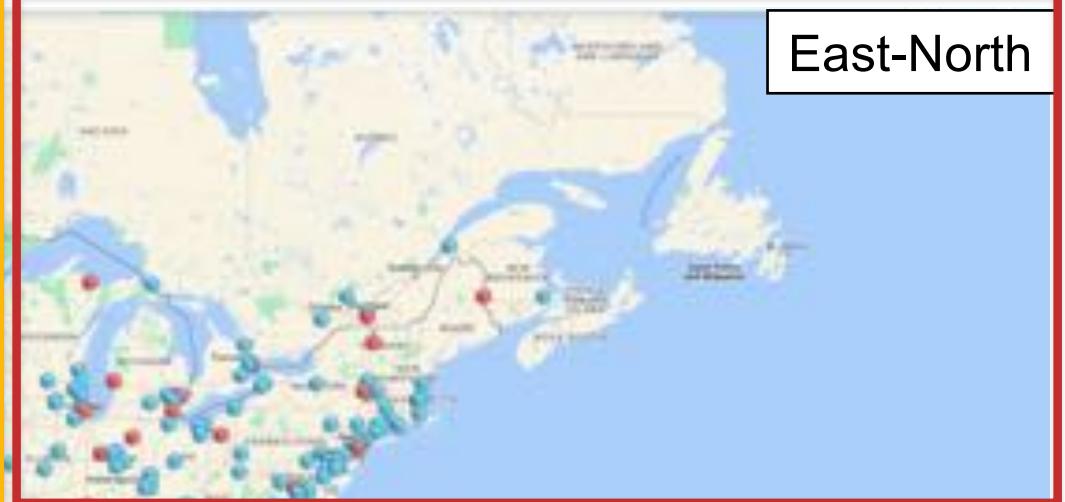
Ideal split will be exactly 50/50 (median east-west coordinate)

Show me the photos!

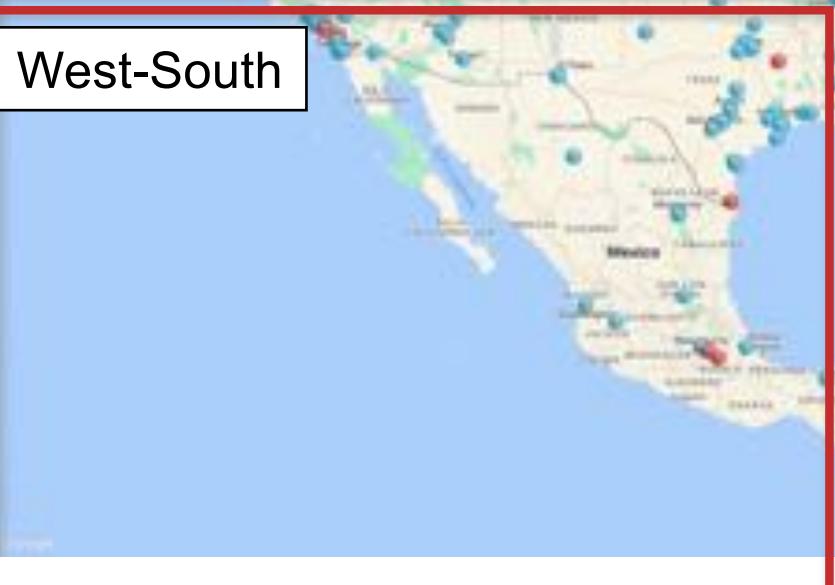
West-North



East-North



West-South



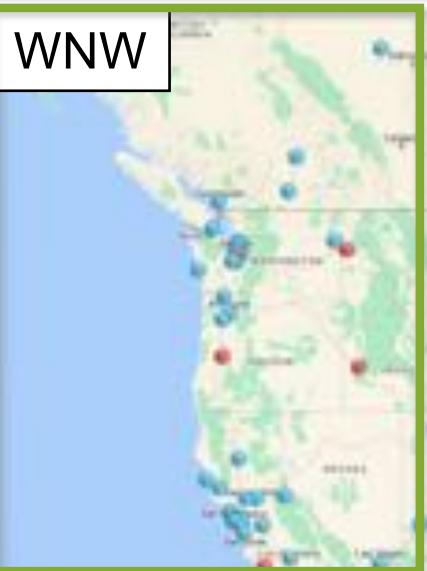
East-South



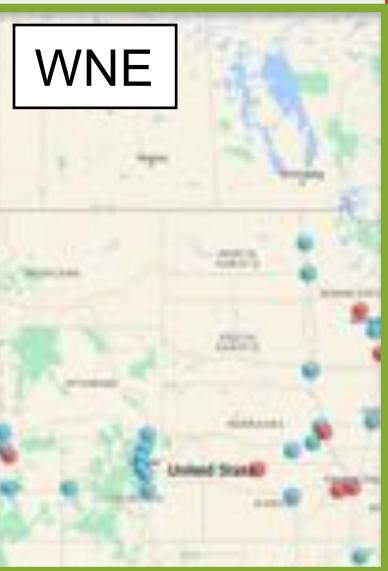
Alternate splits: Each sublist has $N/4$ element & balanced in both dimensions

Show me the photos!

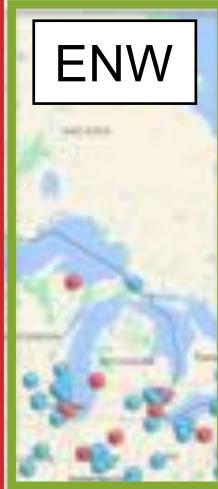
WNW



WNE



ENW



ENE



WSW



WSE



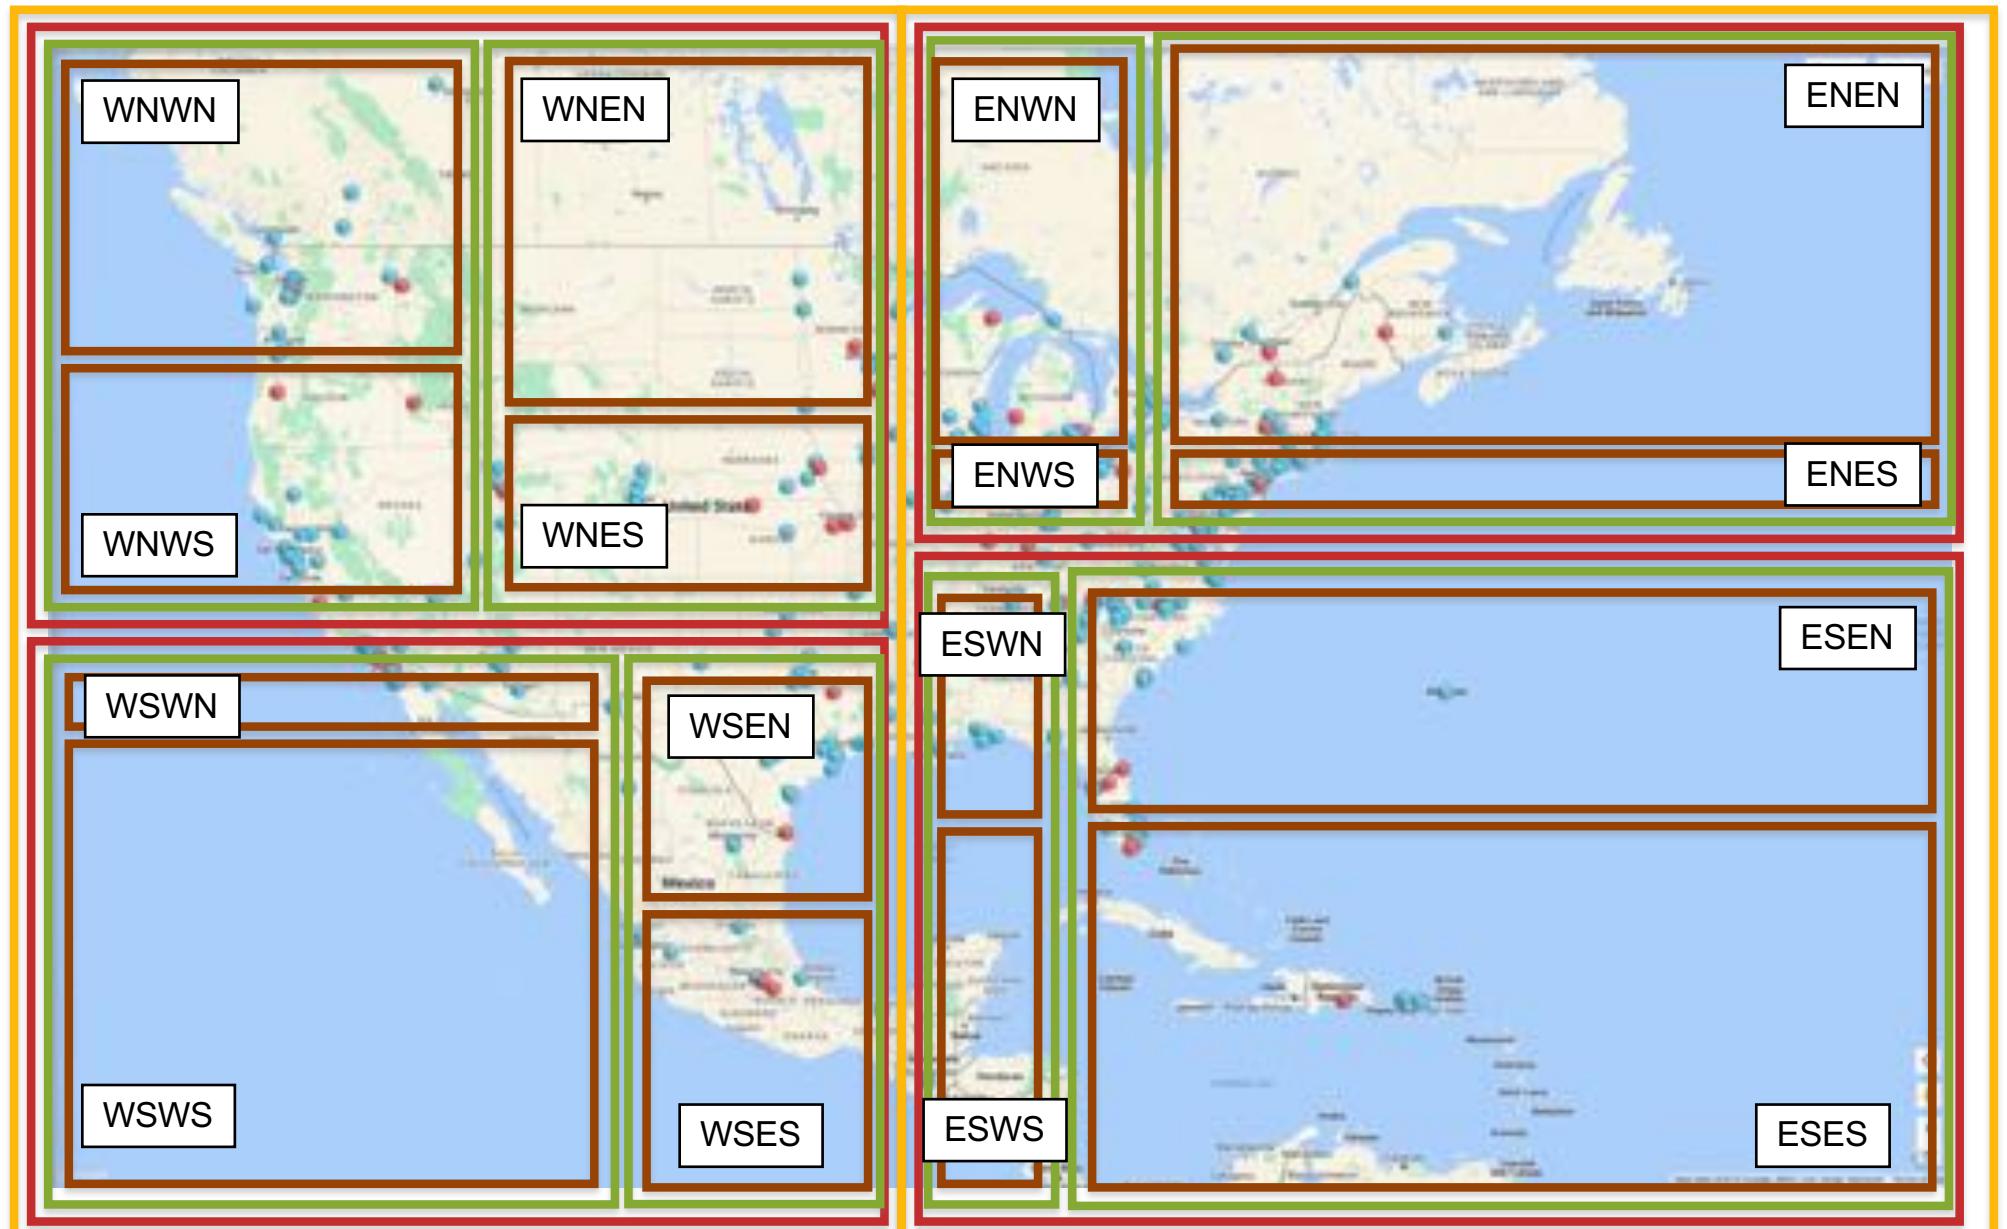
ESW



ESE

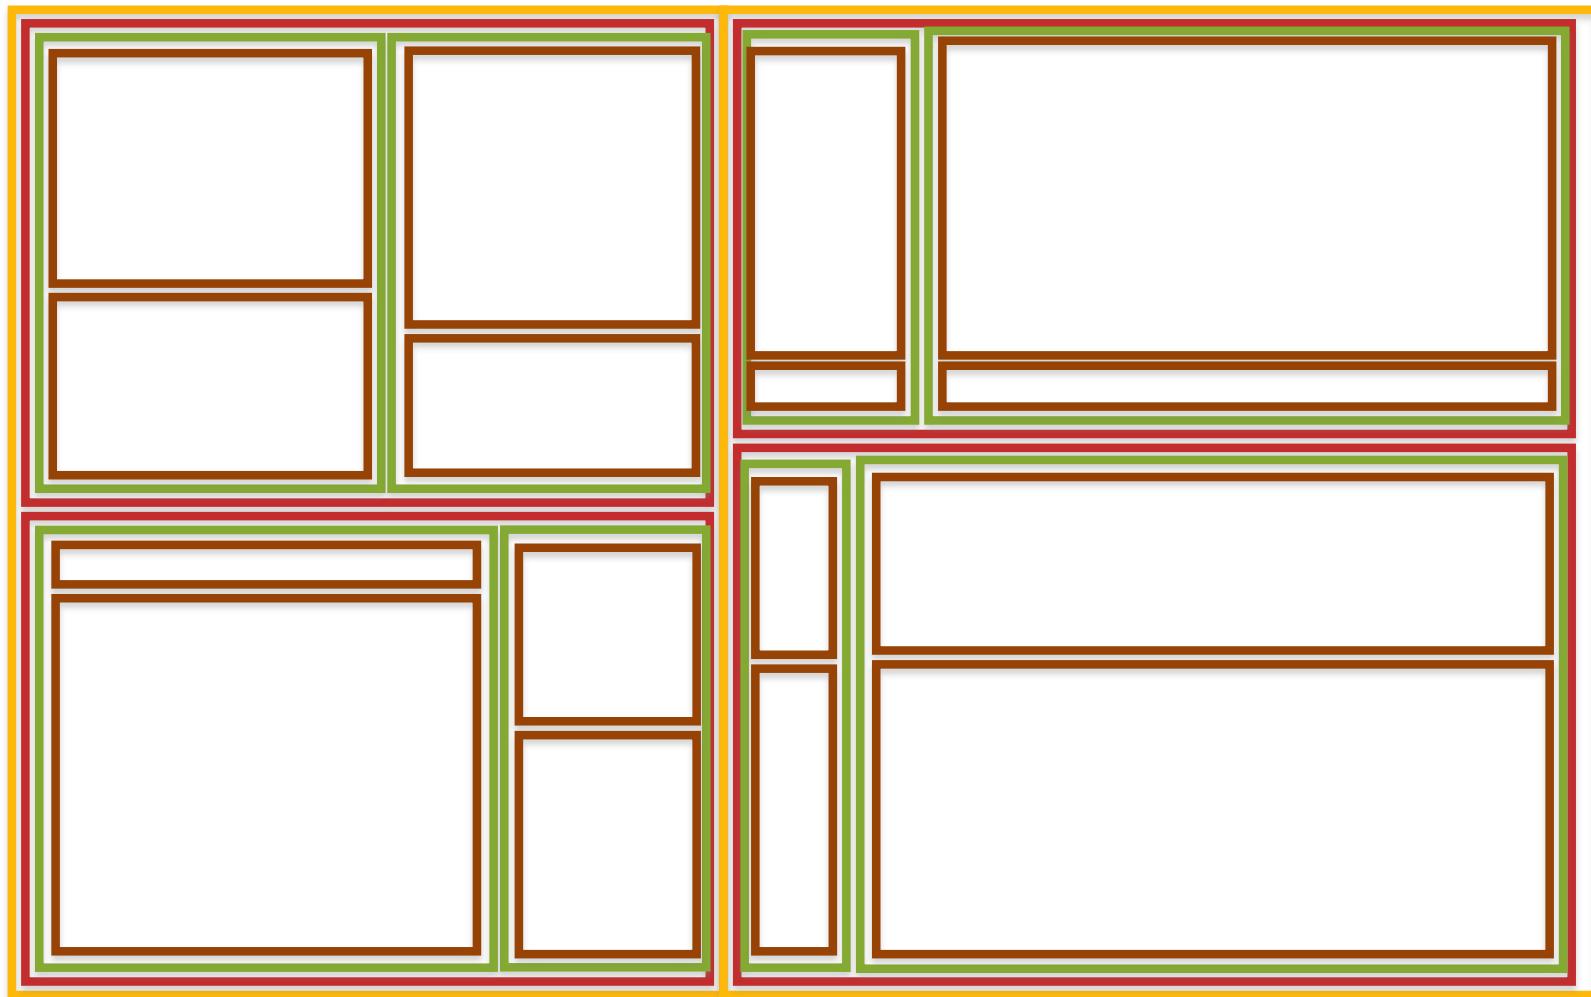
Each sublist has N/8 elements & Balanced in both dimensions

Show me the photos!



Each sublist has $N/16$ elements & Balanced in both dimensions

Advanced Data Structure #1: K-d tree



Balanced Binary Search Tree invented by Jon Louis Bentley in 1975

Generalization of the ubiquitous binary search tree

Very fast to build & search almost any type of spatial data

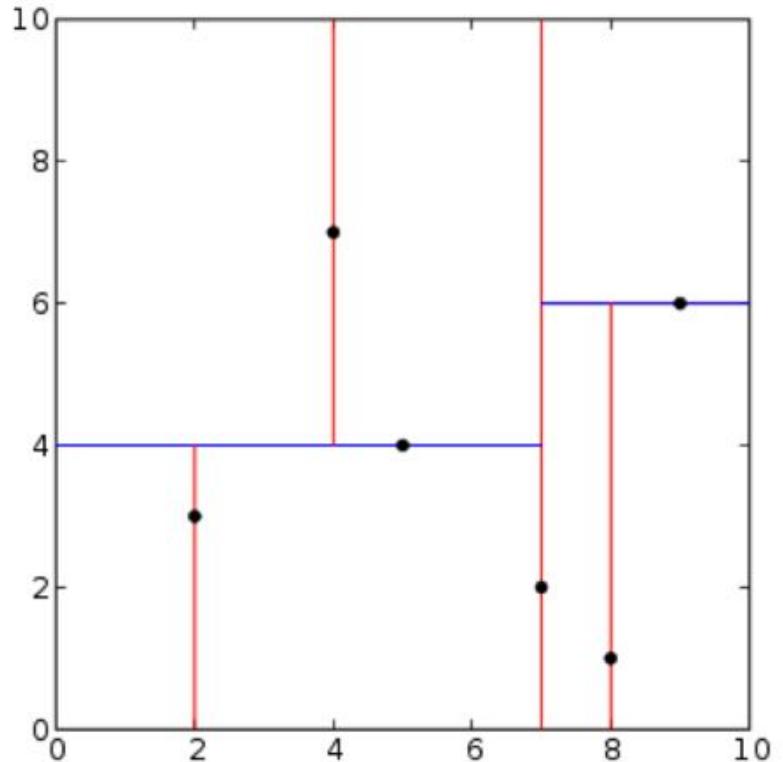
k-d tree pseudocode

```
Photo getNearest(Point myLoc)
{
    // Region class stores partitions & photos
    Region r = allPhotos

    // While more partitions to go
    while (r.numPhotos() > 1)
    {
        // Partition on Lat/Long
        Dimension d = r.splitDim()

        // Check the relevant coordinate
        if (myLoc.getDim(d) <= r.split)
        {
            // branch to the west/south
            r = r.lo()
        }
        else
        {
            // branch to the east/north
            r = r.hi()
        }
    }

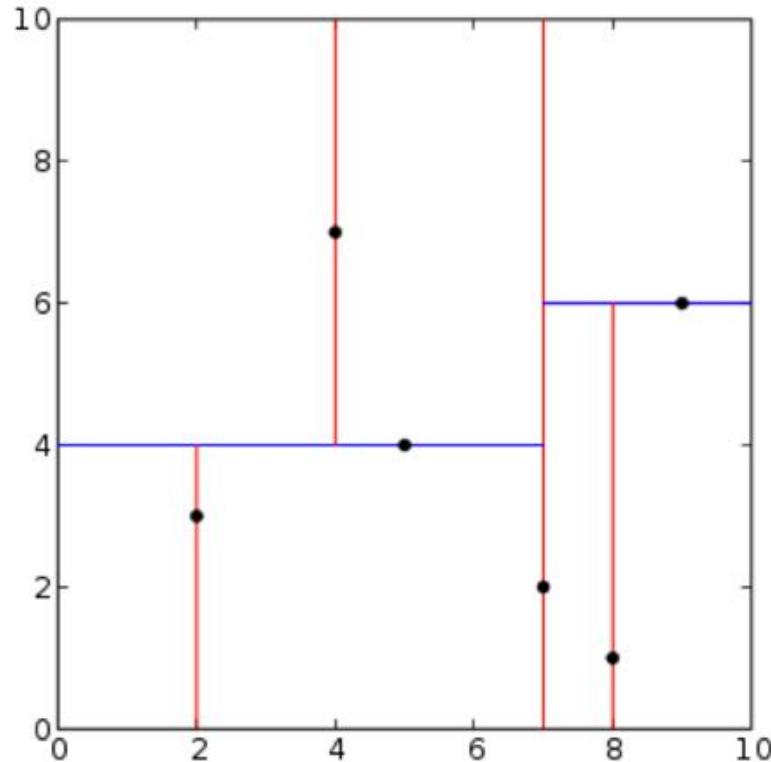
    // just 1 photo, done!
    return r.getPhoto()
}
```



K-d tree data structure to spatially index a large data index the photos

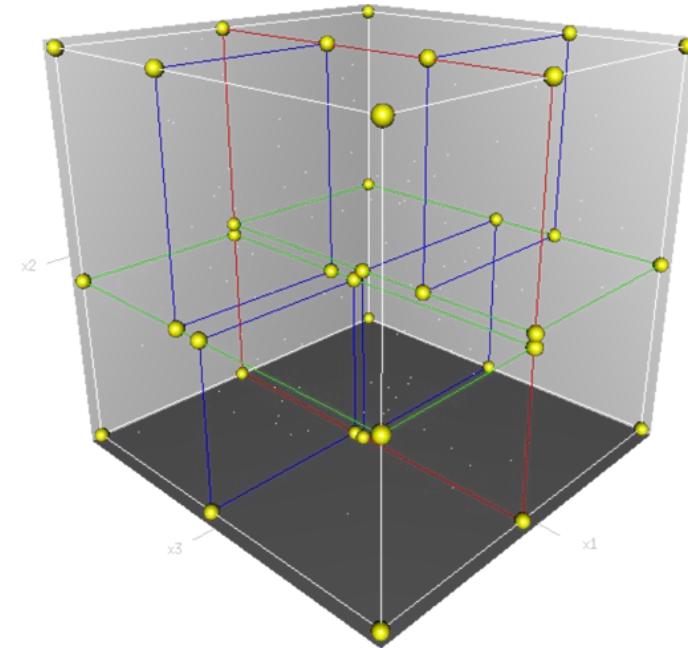
What else might you want to index?

k-d trees in higher dimensions



2d tree:

Alternate left/right, top/bottom



3d tree:

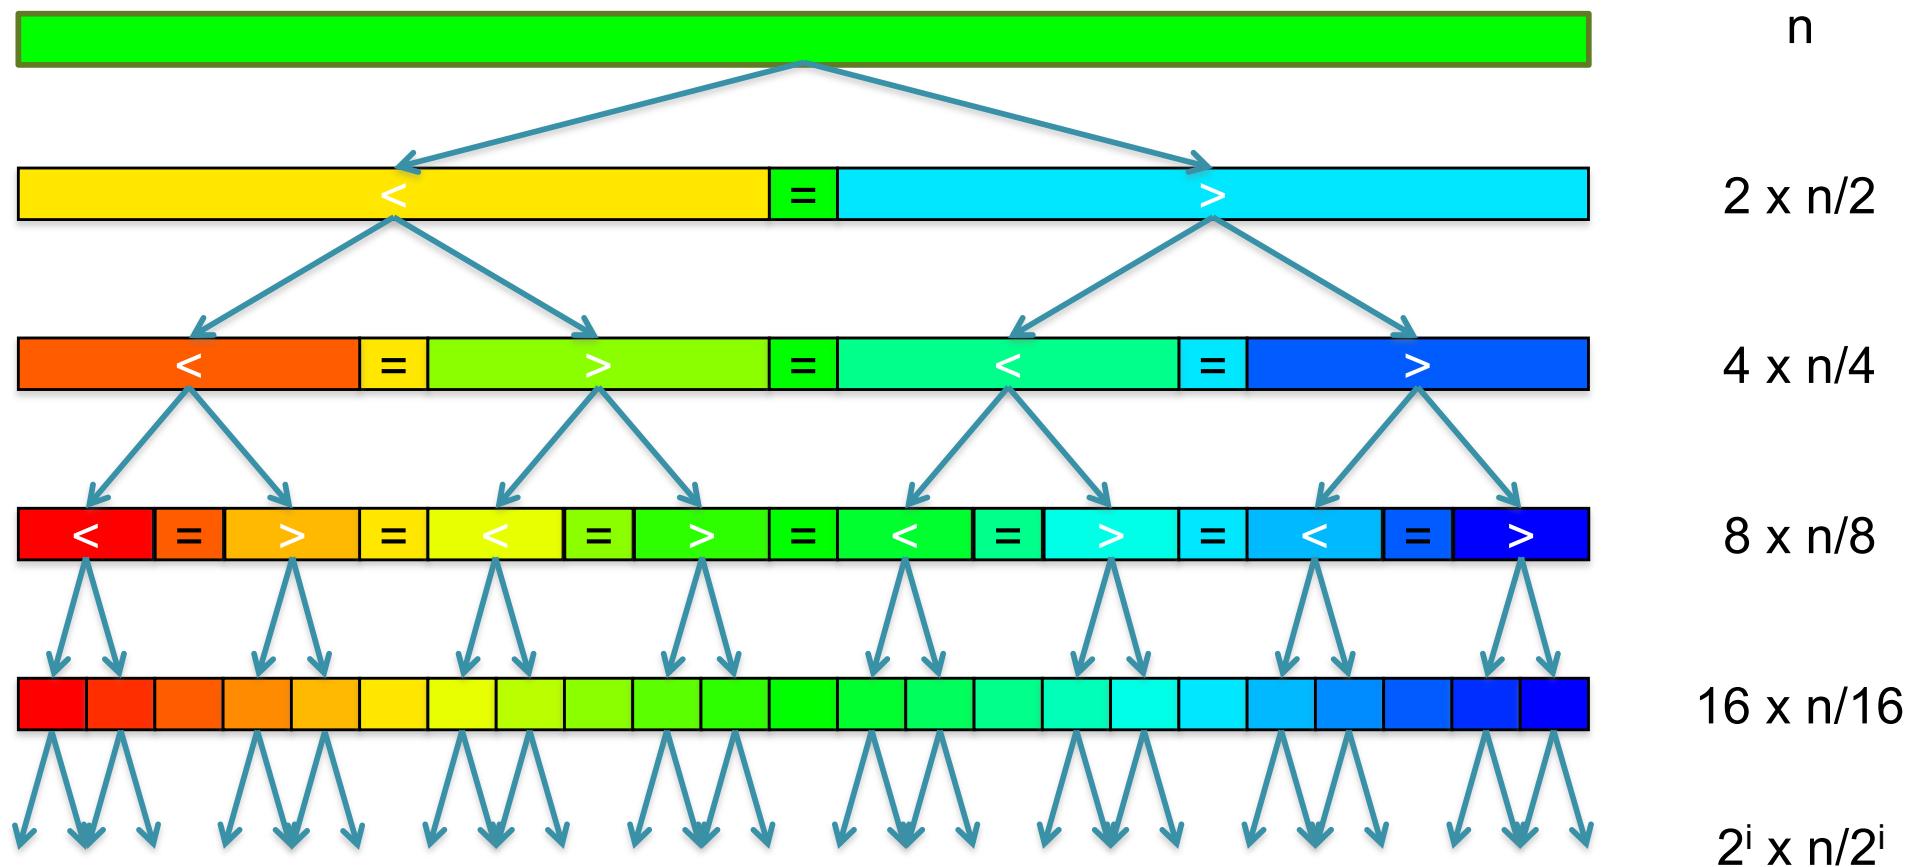
Alternate left/right, top/bottom, up/down

The 'k' in k-d tree emphasizes that it works in any number of dimensions
Just gets a little harder to draw for $k > 3$ ☺

Alternative is to build multiple indices with pointers (URLs) to same set of photos

Divide and Conquer

- Brute force is slow because we have to check every single element
 - How can we split up the unsorted list into independent ranges?
 - Lets recursively split up the elements into greater than/less than range based on the current split line (latitude/longitude)

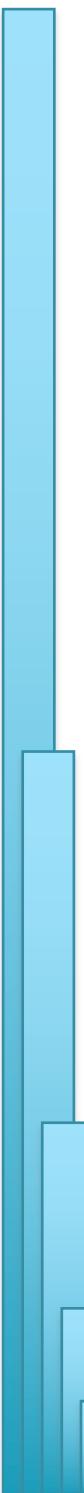


[How many times can we split a list in half?]



Dividing N in half: 20 Billion

- Step 0: 20,000,000,000 possible elements (N)
- Step 1: 10,000,000,000 possible elements ($N/2$)
- Step 2: 5,000,000,000 possible elements ($N/4$)
- ...
- Step X: 1 possible element (N/N)



Dividing N in half: 20 Billion

- Step 0: 20,000,000,000 possible elements ($N/1 = N/2^0$)
- Step 1: 10,000,000,000 possible elements ($N/2 = N/2^1$)
- Step 2: 5,000,000,000 possible elements ($N/4 = N/2^2$)
- ...
- Step X: 1 possible element ($N/N = N/2^X$)

Dividing N in half: 20 Billion

- Step 0: 20,000,000,000 possible elements ($N/1 = N/2^0$)
- Step 1: 10,000,000,000 possible elements ($N/2 = N/2^1$)
- Step 2: 5,000,000,000 possible elements ($N/4 = N/2^2$)
- ...
- Step X: 1 possible element ($N/N = N/2^X$)

Find X such that $2^X \geq N$

Dividing N in half: 20 Billion

- Step 0: 20,000,000,000 possible elements ($N/1 = N/2^0$)
- Step 1: 10,000,000,000 possible elements ($N/2 = N/2^1$)
- Step 2: 5,000,000,000 possible elements ($N/4 = N/2^2$)
- ...
- Step X: 1 possible element ($N/N = N/2^X$)

Find X such that:

$$2^X \geq N$$
$$\lg(2^X) \geq \lg(N)$$
$$X \geq \lg(N)$$

X = ???

Dividing N in half: 20 Billion

- Step 0: 20,000,000,000 possible elements ($N/1 = N/2^0$)
- Step 1: 10,000,000,000 possible elements ($N/2 = N/2^1$)
- Step 2: 5,000,000,000 possible elements ($N/4 = N/2^2$)
- ...
- Step X: 1 possible element ($N/N = N/2^X$)

Find X such that:

$$\begin{aligned}2^X &\geq N \\ \lg(2^X) &\geq \lg(N) \\ X &\geq \lg(N)\end{aligned}$$

X = 35

571.4 million times faster than brute force!

Dividing N in half: 20 TRILLION

- Step 0: 20,000,000,000,000 possible elements ($N/1 = N/2^0$)
- Step 1: 10,000,000,000,000 possible elements ($N/2 = N/2^1$)
- Step 2: 5,000,000,000,000 possible elements ($N/4 = N/2^2$)
- ...
- Step X: 1 possible element ($N/N = N/2^X$)

Find X such that:

$$\begin{aligned}2^X &\geq N \\ \lg(2^X) &\geq \lg(N) \\ X &\geq \lg(N)\end{aligned}$$

X = ???

Dividing N in half: 20 TRILLION

- Step 0: 20,000,000,000,000 possible elements ($N/1 = N/2^0$)
- Step 1: 10,000,000,000,000 possible elements ($N/2 = N/2^1$)
- Step 2: 5,000,000,000,000 possible elements ($N/4 = N/2^2$)
- ...
- Step X: 1 possible element ($N/N = N/2^X$)

Find X such that:

$$\begin{aligned}2^X &\geq N \\ \lg(2^X) &\geq \lg(N) \\ X &\geq \lg(N)\end{aligned}$$

X = 45

571.4 billion times faster than brute force!

Dividing N in half: 20 QUADRILLION

- Step 0: 20,000,000,000,000,000 possible elements ($N/1 = N/2^0$)
- Step 1: 10,000,000,000,000,000 possible elements ($N/2 = N/2^1$)
- Step 2: 5,000,000,000,000,000 possible elements ($N/4 = N/2^2$)
- ...
- Step X: 1 possible element ($N/N = N/2^X$)

Find X such that:

$$2^X \geq N$$
$$\lg(2^X) \geq \lg(N)$$
$$X \geq \lg(N)$$

X = ???

Dividing N in half: 20 QUADRILLION

- Step 0: 20,000,000,000,000,000 possible elements ($N/1 = N/2^0$)
- Step 1: 10,000,000,000,000,000 possible elements ($N/2 = N/2^1$)
- Step 2: 5,000,000,000,000,000 possible elements ($N/4 = N/2^2$)
- ...
- Step X: 1 possible element ($N/N = N/2^X$)

Find X such that:

$$\begin{aligned}2^X &\geq N \\ \lg(2^X) &\geq \lg(N) \\ X &\geq \lg(N)\end{aligned}$$

X = 55

571.4 trillion times faster than brute force!

How much is a zettabyte?

Unit	Size	$\sim 2^x$
Byte	1	2^0
Kilobyte	1,000	2^{10}
Megabyte	1,000,000	2^{20}
Gigabyte	1,000,000,000	2^{30}
Terabyte	1,000,000,000,000	2^{40}
Petabyte	1,000,000,000,000,000	2^{50}
Exabyte	1,000,000,000,000,000,000	2^{60}
Zettabyte	1,000,000,000,000,000,000,000	2^{70}

How much is a zettabyte?

Unit	Size	$\sim 2^x$
Byte	1	2^0
Kilobyte	1,000	2^{10}
Megabyte		2^{20}
Gigabyte		2^{30}
Terabyte	$\lg(X) << 70$	2^{40}
Petabyte		2^{50}
Exabyte	1,000,000,000,000,000,000	2^{60}
Zettabyte	1,000,000,000,000,000,000,000	2^{70}

Next Steps

1. Reflect on the magic and power of log 😊
2. Register on Piazza
3. Set up Dropbox for yourself!
4. Get comfortable with a editor (VI rules!) and the command line



Welcome to CS 600.226

<https://github.com/schatzlab/datastructures2018>

Questions?